

PRODUCTION-GRADE PYTHON · FASTAPI

FastAPI for AI Engineers

From First Endpoint to Production-Scale AI Systems

Covering REST design, Pydantic validation, async patterns, databases, security, LLM & RAG serving, and system design — with 100 interview questions.



AI Engineering Insider

www.aiengineeringinsider.com · aiengineeringinsider.substack.com/subscribe

FastAPI for AI Engineers:

From First Endpoint to Production-Scale AI Systems

Copyright © 2026 AI Engineering Insider. All rights reserved.

No part of this publication may be reproduced or transmitted in any form without prior written permission, except for brief quotations in reviews.

FastAPI is a trademark of Sebastián Ramírez. Python is a trademark of the Python Software Foundation. All other trademarks are the property of their respective owners. The publisher assumes no responsibility for errors or omissions, or for damages resulting from the use of the information contained herein.

First Edition.

Contents

1	Introduction to FastAPI and Modern API Development	9
1.1	What Is FastAPI?	9
1.2	REST API Fundamentals	10
1.3	FastAPI vs Flask vs Django	10
1.4	ASGI vs WSGI: Why the Server Interface Matters	11
1.4.1	<i>The Request Path, End to End</i>	11
1.5	Why FastAPI Dominates AI/ML Serving	11
1.6	Installation and Your First Endpoint	12
1.6.1	<i>Exploring /docs and /redoc</i>	13
1.7	Interview Questions	14
	<i>Q1. What is FastAPI, and what are the two libraries it is built on?</i>	14
	<i>Q2. Explain the difference between WSGI and ASGI. Why can't WSGI support WebSockets?</i>	14
	<i>Q3. When would you choose Django over FastAPI, and vice versa?</i>	14
	<i>Q4. How does FastAPI generate its interactive documentation, and why is this more reliable than hand-written docs?</i>	15
	<i>Q5. A <code>sync def predict()</code> spends 200ms in-process; an <code>async def fetch()</code> awaits a 200ms upstream call. How does FastAPI schedule each, and what are the throughput implications?</i>	15
	<i>Q6. What is OpenAPI, and how does it differ from Swagger UI?</i>	15
	<i>Q7. Why are Python type hints central to FastAPI rather than just stylistic?</i>	15
	<i>Q8. What does "statelessness" mean in REST, and what does it buy you operationally?</i>	16
	<i>Q9. What is Uvicorn, and why does FastAPI need a separate server at all?</i>	16
	<i>Q10. Your team's hand-maintained API docs keep drifting from reality and integration partners complain. How would FastAPI change this workflow?</i>	16
2	Python Typing, Pydantic, and Data Validation	17
2.1	Python Type Hints: The Substrate	17
2.2	BaseModel: Declaring Schemas	17
2.3	Request and Response Schemas in FastAPI	18
2.4	Nested Models and Field Validation	19
2.5	Pydantic v2 Essentials	20
2.6	Preventing Sensitive Data Leaks	20
2.7	Interview Questions	22
	<i>Q1. Python type hints are not enforced by the interpreter. How does Pydantic make them enforceable, and what is the cost?</i>	22

Q2. Why should request, storage, and response schemas be separate models even though it triples the class count?	22
Q3. Explain Pydantic v2's lax vs strict modes. When is coercion dangerous?	22
Q4. An endpoint returns the SQLAlchemy User object directly and the response contains hashed_password. Walk through the layers of defense you would add.	22
Q5. How do nested model validation errors surface to the client, and why is the error structure important for API consumers?	23
Q6. What is the difference between a required field, an optional field, and a nullable field in Pydantic?	23
Q7. What does Field(default_factory=list) solve, and why not write tags: list[str] = []?	23
Q8. How do Pydantic Field constraints interact with the OpenAPI docs?	24
Q9. When would you use @model_validator instead of @field_validator?	24
Q10. Your API's p99 latency budget is 50ms and payloads are 2MB nested JSON. How do you keep validation from blowing the budget?	24
3 Building Core API Endpoints	25
3.1 Verbs, Semantics, and Idempotency	25
3.2 Path Parameters, Query Parameters, and Bodies	25
3.3 Status Codes That Mean What They Say	27
3.4 Pagination, Filtering, and Sorting	27
3.5 API Versioning with /api/v1	28
3.6 Interview Questions	29
Q1. Explain idempotency for each HTTP verb, and how you would make POST safe to retry.	29
Q2. Offset vs cursor pagination: trade-offs, failure modes, and when each is correct.	29
Q3. A client PATCHes {"description": null}. How do you distinguish "clear this field" from "field not provided," and how does Pydantic support it?	29
Q4. Walk through correct status-code choices for: duplicate signup email, expired token, valid token but wrong role, malformed JSON, schema-valid but business-rule-violating payload, and a deleted resource.	30
Q5. How would you evolve an API from v1 to v2 without breaking existing consumers? Describe the full lifecycle.	30
Q6. Why must list-endpoint limit parameters have a server-enforced maximum?	30
Q7. What belongs in a path parameter vs a query parameter, and why does the distinction matter?	30
Q8. What is the practical difference between PUT and PATCH for an update endpoint?	31
Q9. How do tags, summaries, and descriptions on endpoints affect more than just aesthetics?	31
Q10. Your list endpoint is p95-slow only for a handful of consumers. Investigation shows they paginate to page 4,000. What are your options?	31
4 Dependency Injection and Application Structure	32
4.1 The Dependency Injection Concept	32
4.1.1 Router-Level and App-Level Dependencies	33
4.2 Routers and Project Structure	33
4.3 Clean Architecture: Router, Service, Repository	34
4.4 Interview Questions	36
Q1. Explain how FastAPI resolves Depends() for a request, including chaining and caching.	36
Q2. Why do yield-based dependencies matter for database session management? What happens on an unhandled exception?	36

Q3. Where should business logic live, and how do you keep HTTP concerns out of it? What goes wrong otherwise?	37
Q4. The repository pattern looks like indirection for its own sake — “my service could just call SQLAlchemy.” Defend it, and state when you would skip it.	37
Q5. Compare router-level dependencies, app-level dependencies, and middleware for enforcing authentication.	37
Q6. What is a composition root, and why should <code>Depends()</code> wiring be centralized?	38
Q7. How does <code>dependency_overrides</code> relate to this chapter's architecture?	38
Q8. When does a class-based dependency beat a function-based one?	38
Q9. Five dependencies in one request all declare <code>Depends(get_settings)</code> . How many times does <code>get_settings</code> execute, and how do you reduce that to once per process?	39
Q10. Your codebase constructs services inline inside every handler. List the concrete problems this causes and outline the migration to <code>Depends</code> -based wiring.	39
5 Database Integration with SQLAlchemy, SQLModel, and Alembic	40
5.1 SQLite for Development, PostgreSQL for Production	40
5.2 Engine, Sessions, and the Unit of Work	40
5.2.1 CRUD Against Models	41
5.3 SQLModel: One Class, Two Roles	42
5.4 Relationships, Lazy Loading, and the N+1 Problem	42
5.5 Alembic Migrations	42
5.6 Interview Questions	44
Q1. Explain the N+1 query problem, how to detect it in a FastAPI service, and the trade-offs between <code>selectinload</code> and <code>joinedload</code> .	44
Q2. Walk through connection pool sizing for 8 pods × 4 workers against Postgres <code>max_connections=200</code> . What breaks first if you get it wrong, and where does PgBouncer fit?	44
Q3. Design the transaction boundary for “create order, decrement inventory, enqueue confirmation email.” Which parts belong in the DB transaction and which must not be in it?	45
Q4. Your team deploys with zero downtime. Describe how to rename a column safely, and why <code>alembic revision -autogenerate</code> alone is insufficient.	45
Q5. Compare SQLModel and plain SQLAlchemy + separate Pydantic schemas for a growing service.	46
Q6. What does the session-per-request pattern guarantee, and what goes wrong with a module-level global session?	46
Q7. Why is <code>pool_pre_ping</code> (or an equivalent) needed in cloud deployments?	46
Q8. What is the difference between <code>flush()</code> and <code>commit()</code> , and when do you need flush explicitly?	46
Q9. Reads are 20× writes and reporting queries are degrading checkout latency. Options?	47
Q10. How would you make repository-layer code testable without mocking SQLAlchemy internals?	47
6 Authentication, Authorization, and API Security	48
6.1 Authentication vs Authorization	48
6.2 Password Hashing	48
6.3 JWT Access Tokens and the Current-User Dependency	49
6.4 API Keys, CORS, Rate Limiting, and Headers	50
6.5 OWASP API Security: Where Breaches Actually Happen	51
6.6 Interview Questions	52
Q1. Why must JWT verification pin the algorithm, and what is the <code>alg=none</code> attack?	52

Q2. Design a revocation story for stateless JWTs. Walk through the trade-offs.	52
Q3. What is BOLA, why does it top the OWASP API list, and how do you make it structurally impossible rather than reviewer-dependent?	53
Q4. Compare storing web-client tokens in <code>localStorage</code> vs <code>httpOnly</code> cookies, including the CSRF implications.	53
Q5. Why <code>bcrypt/argon2</code> instead of SHA-256 for passwords, and how do you choose/maintain cost parameters in production?	53
Q6. A login endpoint returns “email not found” vs “wrong password.” What is wrong, and what else leaks account existence?	53
Q7. What does CORS actually protect, and why is it not an access-control mechanism?	54
Q8. Where should rate limiting live — gateway, middleware, or dependency — and what does Redis add?	54
Q9. Your JWT secret leaked in a public repo. Walk through the incident response.	54
Q10. When do API keys beat OAuth2/JWT for authentication, and what hygiene do they require?	55
7 Testing FastAPI Applications	56
7.1 Why API Testing Is Different	56
7.2 <code>pytest</code> and <code>TestClient</code>	56
7.3 Testing CRUD, Validation, and Contracts	57
7.4 Testing Authentication and Authorization	58
7.4.1 Mocking External Dependencies	58
7.5 Test Databases and CI Workflow	59
7.6 Code Coverage, Honestly	59
7.7 Interview Questions	61
Q1. How do FastAPI dependency overrides work, and why are they superior to monkeypatching for API tests?	61
Q2. Design the test-database strategy for a team using Postgres in production. Defend SQLite-in-memory against “test what you fly.”	61
Q3. What belongs in a unit test vs an integration test for a FastAPI endpoint, and how does the Chapter-4 architecture determine this?	61
Q4. How would you test that an endpoint never leaks sensitive fields, in a way that survives future refactors?	62
Q5. Your CI suite has a 3% flake rate across 2,000 tests. Quantify the damage and lay out a remediation plan.	62
Q6. Why does <code>TestClient</code> not require a running server, and what are the limits of in-process testing?	62
Q7. How do you test code that depends on the current time (token expiry, rate limits)?	63
Q8. What does <code>with TestClient(app) as client</code> do that bare <code>TestClient(app)</code> does not?	63
Q9. Why is asserting on the 422 error body (not just the status code) worth the test brittleness?	63
Q10. Your coverage is 95% but production incidents keep coming from “tested” code. Diagnose.	63
8 Async Programming, Background Tasks, and External Services	64
8.1 The Event Loop Mental Model	64
8.2 Calling External APIs with <code>httpx</code>	65
8.3 <code>BackgroundTasks</code> vs Celery: Two Tiers of Deferred Work	65
8.4 Webhooks, Uploads, Streaming, and WebSockets	66
8.5 Interview Questions	68
Q1. Explain precisely what happens when an <code>async def</code> FastAPI handler makes a blocking call. Why does it degrade endpoints that share nothing with it?	68

Q2. Design a retry policy for calls to a flaky upstream. What can naive retries do to a degraded service, and what are the safeguards?	68
Q3. BackgroundTasks vs Celery: give the decision criteria and the failure modes of choosing wrong.	68
Q4. How do you receive webhooks safely? Cover authenticity, replay, ordering, and the timeout contract.	69
Q5. When does WebSocket beat SSE/polling, and what operational complexity does it introduce?	69
Q6. Why must the shared <code>httpx.AsyncClient</code> live in the lifespan handler rather than being created per request?	70
Q7. What is the difference between <code>StreamingResponse</code> and building the full response in memory, and when is streaming mandatory?	70
Q8. How do you enforce an upload size limit robustly, and why is checking <code>Content-Length</code> insufficient?	70
Q9. Your Celery queue depth is growing without bound. Walk through the diagnosis and the levers.	70
Q10. Explain backpressure in the context of an async API. Where does it come from and why is unbounded concurrency dangerous?	71
9 FastAPI for AI, ML, RAG, and LLM Applications	72
9.1 Serving ML Models: The Shape of the Problem	72
9.1.1 Batch Inference	73
9.2 Wrapping LLMs: Gateway Patterns and Streaming	73
9.3 Embeddings, Vector Databases, and RAG	74
9.4 Guardrails: Validating Inputs and Outputs of a Stochastic System	74
9.5 Versioning, Monitoring, and Humans in the Loop	75
9.6 Interview Questions	77
Q1. Why must ML models load in the lifespan handler, and what goes wrong with lazy per-request or import-time loading?	77
Q2. Design the streaming architecture for an LLM chat endpoint end to end — protocol choice, disconnects, metering, and what changes at the infrastructure layer.	77
Q3. A RAG system gives confidently wrong answers. Walk through your diagnosis across the pipeline.	77
Q4. How do you make a stochastic LLM safe to use as a typed component in a larger system?	78
Q5. Compare pgvector against a dedicated vector database for a RAG product, and give the migration trigger points.	78
Q6. Why does an embedding API need the model version in its response, and what is the operational consequence of upgrading the embedding model?	78
Q7. What is prompt injection, and which defenses actually work for an LLM endpoint that processes user documents?	79
Q8. Your LLM bill doubled month-over-month with flat traffic. Where do you look?	79
Q9. When does batch inference beat real-time, and how do you expose each through an API?	79
Q10. Design the human-in-the-loop component for an AI document-extraction service. What makes it an engineering problem rather than a UI problem?	80
10 Production Deployment, Scaling, and Observability	81
10.1 Servers and Workers: Uvicorn and Gunicorn	81
10.2 Dockerizing FastAPI	81
10.3 Configuration, Cloud Targets, and Kubernetes	82
10.4 Observability: Logs, Metrics, Traces	83

10.5 Caching, Performance Tuning, and Load Testing	84
10.6 The Production Checklist	85
10.7 Interview Questions	86
Q1. Why do liveness and readiness probes need different implementations, and what failure mode does conflating them create?	86
Q2. Design a zero-downtime deploy for a FastAPI service with database migrations and 30-second-lived in-flight requests.	86
Q3. Your p99 latency tripled but p50 is unchanged. Walk through your diagnosis using the three observability pillars.	86
Q4. Cache-aside with TTL: walk through the stampede problem and its solutions, and explain why invalidation is harder than caching.	87
Q5. How would you load test this book's RAG service meaningfully? What do naive load tests get wrong?	87
Q6. Gunicorn with Uvicorn workers vs plain Uvicorn <code>-workers</code> vs one-process-per-pod on Kubernetes: how do you choose?	87
Q7. What belongs in structured logs for an API, and what must never appear?	88
Q8. Explain SLO-based alerting and why paging on "CPU above 80%" is an anti-pattern.	88
Q9. Your service must survive its primary LLM provider's outage. What does graceful degradation look like? . . .	88
Q10. A pod is OOM-killed every few hours under steady traffic. Walk through the investigation.	89

Introduction to FastAPI and Modern API Development

This chapter establishes the foundation for the entire book: what FastAPI is, why it dominates modern Python API development, how it compares to Flask and Django, what ASGI means and why it matters, and why FastAPI has become the de facto serving layer for AI and ML systems. By the end you will have a running API with automatic interactive documentation.

1.1 What Is FastAPI?

FastAPI is a modern Python web framework for building HTTP APIs, created by Sebastián Ramírez and first released in 2018. Its design rests on one deceptively simple idea: *the Python type hints you already write should do the work*. When you annotate a function parameter as `int`, FastAPI uses that single annotation to parse the value from the request, validate it, convert it, document it in the OpenAPI schema, and render it in the interactive docs. One declaration, five behaviors.

FastAPI is deliberately a thin composition of two excellent libraries rather than a monolith. **Starlette** provides the ASGI web toolkit: routing, middleware, WebSockets, background tasks, and the request/response cycle. **Pydantic** provides data parsing and validation driven by type annotations. FastAPI's contribution is the dependency injection system, the OpenAPI integration, and the developer ergonomics that glue them together. This layered design matters in practice: when you debug a FastAPI application, you are often actually reading Starlette or Pydantic source, and understanding the boundary between the three layers is what separates a user of the framework from an engineer who can reason about it.

• CORE CONCEPT — Declarative request contracts

The central mental model of FastAPI: an endpoint signature is a *contract*. Parameters declare what the request must contain; the return annotation (or `response_model`) declares what the response will contain. The framework enforces the contract at runtime and publishes it as machine-readable OpenAPI. Everything else in this book — validation, documentation, security, testing — builds on this single idea.

1.2 REST API Fundamentals

Before comparing frameworks, we need shared vocabulary. A **REST API** (REpresentational State Transfer) exposes *resources* (users, orders, model predictions) at URLs, manipulated through standard HTTP verbs: `GET` reads, `POST` creates, `PUT` replaces, `PATCH` partially updates, and `DELETE` removes. Responses carry standardized status codes — `2xx` success, `4xx` client error, `5xx` server error — and typically JSON bodies.

REST is an architectural style, not a protocol. Its core constraints are statelessness (every request carries all the context the server needs, which is what makes horizontal scaling trivial), a uniform interface, and cacheability. Real-world “REST” APIs are usually pragmatic approximations, and that is fine; what matters is consistency: predictable URLs, correct verbs, correct status codes, and stable schemas.

OpenAPI (formerly Swagger) is the industry-standard, machine-readable description of an HTTP API: every path, parameter, request body, response schema, and security scheme, expressed as JSON. From an OpenAPI document you can generate client SDKs, server stubs, contract tests, and interactive documentation. Most frameworks treat OpenAPI as an afterthought maintained by hand; FastAPI generates it automatically from your code, so the documentation can never drift out of sync with the implementation.

1.3 FastAPI vs Flask vs Django

Dimension	FastAPI	Flask	Django (+DRF)
Server interface	ASGI (async-native)	WSGI (sync; async via extensions)	WSGI + ASGI (hybrid)
Validation	Built-in via Pydantic	Manual / extensions	DRF serializers
OpenAPI docs	Automatic	Extensions (flasgger)	Extensions (drf-spectacular)
Type-hint driven	Core design principle	Not used	Not used
Batteries included	Minimal core	Minimal core	ORM, admin, auth, templates
Concurrency model	Event loop + threadpool	Worker processes/threads	Worker processes/threads
Sweet spot	APIs, microservices, ML serving	Small apps, quick prototypes	Full-stack web apps, CMS, admin-heavy products

Table 1.1: Framework comparison at a glance.

Flask (2010) is a WSGI microframework: explicit, unopinionated, and tiny. Its weakness for API work is that everything FastAPI gives you for free — validation, serialization, documentation, async concurrency — must be assembled from third-party extensions, each with its own conventions.

Django (2005) is a full-stack framework whose ORM, admin interface, and authentication system remain best-in-class. With Django REST Framework it builds excellent APIs, but the request-validation story (serializers defined separately from views) involves more ceremony, and async support was retrofitted onto a fundamentally synchronous core.

FastAPI was designed API-first in the era of type hints and `asyncio`. Choose it when the product is the API: microservices, mobile/SPA backends, and — the focus of this book — AI/ML inference services. Choose Django when you need its admin and ORM ecosystem for a content-heavy product; choose Flask when you want a minimal sync app and full manual control.

1.4 ASGI vs WSGI: Why the Server Interface Matters

WSGI (Web Server Gateway Interface, 2003) defines a synchronous contract: the server calls your application with a request, and your code blocks until it returns a response. One request occupies one worker thread/process for its entire duration. If your handler spends 2 seconds waiting on a database or an upstream LLM API, that worker does *nothing* for 2 seconds.

ASGI (Asynchronous Server Gateway Interface) replaces the blocking call with an async protocol of `await`-able events. A single event loop can interleave thousands of in-flight requests: while one request awaits a database row, the loop services others. ASGI also supports protocols WSGI structurally cannot: WebSockets, server-sent events, and HTTP/2 streaming — all essential for streaming LLM token responses, as we will see in Chapter 9.

The performance win of ASGI is specifically for **I/O-bound** workloads — waiting on databases, caches, and upstream APIs. For CPU-bound work (in-process model inference, image processing), the event loop does not help and can actively hurt if you block it. Chapter 8 treats this distinction in depth; it is the single most common FastAPI production mistake.

1.4.1 The Request Path, End to End

Figure 1.1 shows the architecture you are adopting when you deploy FastAPI: clients reach a load balancer, which fans out to Uvicorn (an ASGI server) processes, each running the FastAPI application built on Starlette and Pydantic, which in turn talks to databases, caches, and model backends.




Figure 1.1: Production FastAPI request path: clients → load balancer → Uvicorn (ASGI) workers → Starlette routing → Pydantic validation → your handlers → data stores.

1.5 Why FastAPI Dominates AI/ML Serving

Nearly every major ML serving stack has converged on FastAPI or its underpinnings: it is the default scaffolding for inference endpoints, LLM gateways, and RAG backends. Five properties explain this:

1. **Python-native.** The models are in Python (PyTorch, transformers, scikit-learn); serving them from the same process eliminates serialization boundaries and language bridges.
2. **Async I/O fits LLM workloads.** Calling a hosted LLM API is a multi-second I/O wait. An async framework holds thousands of these concurrent calls open on a handful of workers.
3. **Streaming support.** Token-by-token streaming over SSE or WebSockets requires ASGI; WSGI frameworks cannot do it cleanly.

4. **Pydantic as a guardrail layer.** Validating prompts in and structured model output out is exactly Pydantic's job; the ecosystem (OpenAI SDK, LangChain, vLLM, Ray Serve) already speaks Pydantic.
5. **Automatic OpenAPI.** ML teams ship APIs consumed by other teams; self-updating interactive docs remove an entire class of integration friction.

1.6 Installation and Your First Endpoint

FastAPI requires Python 3.8+ (3.11+ recommended for meaningful asyncio and interpreter speedups). Install the framework and the Uvicorn server:

```
1 python -m venv .venv && source .venv/bin/activate
2 pip install "fastapi[standard]" # includes uvicorn, httpx, etc.
```

Listing 1.1: Installing FastAPI with all standard extras.

Now the canonical first application — but slightly beyond hello-world, so that every FastAPI superpower is already visible:

```
1 from fastapi import FastAPI
2 from pydantic import BaseModel
3
4 app = FastAPI(
5     title="Model Serving API",
6     version="1.0.0",
7     description="First steps toward a production AI service.",
8 )
9
10 class PredictionRequest(BaseModel):
11     text: str
12     max_tokens: int = 64
13
14 class PredictionResponse(BaseModel):
15     label: str
16     confidence: float
17
18 @app.get("/health")
19 def health() -> dict:
20     return {"status": "ok"}
21
22 @app.get("/items/{item_id}")
23 def read_item(item_id: int, verbose: bool = False) -> dict:
24     # item_id parsed+validated from the path; verbose from ?verbose=
25     return {"item_id": item_id, "verbose": verbose}
26
27 @app.post("/predict", response_model=PredictionResponse)
28 def predict(req: PredictionRequest) -> PredictionResponse:
29     # req is already parsed, validated, and typed.
30     return PredictionResponse(label="positive", confidence=0.97)
```

Listing 1.2: main.py — first endpoints with typed parameters.

Run it with hot reload during development:

```
uvicorn main:app --reload --port 8000
```

Send `GET /items/abc` and FastAPI responds with a structured 422 error explaining that `item_id` must be an integer — you wrote zero validation code.

1.6.1 Exploring /docs and /redoc

Open `http://localhost:8000/docs`: **Swagger UI** renders every endpoint with its schemas and lets you execute requests in the browser. `/redoc` serves the same OpenAPI document in **ReDoc**, a three-panel reference layout better suited to publishing. The raw contract lives at `/openapi.json` — this is the artifact you feed to client generators and contract-testing tools in CI.

PRODUCTION CASE STUDY — Netflix — internal crisis-management APIs on FastAPI

Netflix's Dispatch platform (open-sourced in 2020) coordinates incident response across Slack, Jira, and Google Workspace, and is built on FastAPI with Pydantic models as the single source of truth for every integration payload. The team cited two decisive factors: automatic OpenAPI generation meant internal consumers always had accurate docs without anyone maintaining them, and async handlers let one modest deployment fan out to many third-party APIs concurrently during an incident — exactly when latency matters most. The pattern to copy: *when your service is mostly orchestrating other services over HTTP, an async-native framework converts waiting time into throughput.*

COST MODEL & METRICS — Worker economics: WSGI vs ASGI

Suppose each request spends T_{io} seconds waiting on I/O and T_{cpu} seconds computing. A synchronous worker sustains

$$R_{sync} = \frac{1}{T_{io} + T_{cpu}} \text{ req/s per worker,}$$

while an async worker (one event loop, C in-flight requests) sustains approximately

$$R_{async} \approx \min\left(\frac{1}{T_{cpu}}, \frac{C}{T_{io} + T_{cpu}}\right).$$

With $T_{io} = 300$ ms, $T_{cpu} = 5$ ms: a sync worker delivers ≈ 3.3 req/s; an async worker approaches $1/0.005 = 200$ req/s. To serve 1,000 req/s you need ~ 300 sync workers versus ~ 5 – 8 async workers. At a typical cloud price of \$0.04/vCPU-hour, that is roughly \$8,600/month versus \$230/month — a 97% infrastructure cost reduction for I/O-heavy traffic. This formula recurs throughout the book; internalize it.

SYSTEM DESIGN SCENARIO — Design a public API for a sentiment-analysis model

Requirements: 500 req/s peak, p95 latency < 150 ms, model inference takes 20 ms CPU, callers are external customers.

Reference approach: (1) FastAPI behind a load balancer with 4–6 Uvicorn workers per node; CPU-bound inference means worker count $\approx$ core count, and you scale horizontally on CPU utilization. (2) Pydantic request model caps input length (defense and cost control). (3) `response_model` hides internal fields. (4) `/health` endpoint for the balancer; `/openapi.json` drives the public docs portal. (5) Rate limiting at the gateway (Chapter 6) and Prometheus metrics (Chapter 10). Capacity check: $500 \times 0.020\text{ s} = 10$ busy cores at peak; provision ~ 16 cores across 2+ nodes for headroom and redundancy.

1.7 Interview Questions

Q1. What is FastAPI, and what are the two libraries it is built on?

Answer. FastAPI is an ASGI Python framework for building typed HTTP APIs. It composes **Starlette** (ASGI toolkit: routing, middleware, WebSockets, background tasks) with **Pydantic** (type-annotation-driven parsing and validation), adding its own dependency injection system and automatic OpenAPI generation. A strong answer notes the practical consequence: FastAPI bugs and behaviors often live in Starlette or Pydantic, so debugging requires knowing which layer owns which responsibility — e.g. middleware ordering is Starlette, validation errors are Pydantic, `Depends()` resolution is FastAPI itself.

Q2. Explain the difference between WSGI and ASGI. Why can't WSGI support WebSockets?

Answer. WSGI defines a synchronous call: `app(environ, start_response)` returns a complete response; one worker is occupied per request for its full duration. ASGI defines an async coroutine `app(scope, receive, send)` exchanging event messages, so a single event loop multiplexes many requests. WSGI structurally assumes a *request-then-response* lifecycle that terminates; WebSockets require a long-lived, bidirectional channel where either side sends at any time — there is no place in the WSGI contract for the server to deliver later messages to the handler or for the handler to push multiple frames. ASGI's `receive/send` message pairs model exactly that.

Q3. When would you choose Django over FastAPI, and vice versa?

Answer. Choose Django when the product needs what Django uniquely ships: the admin interface (a free internal CRUD UI), the mature ORM with migrations, built-in auth/sessions, and a large plugin ecosystem — typical for content platforms and admin-heavy business

systems. Choose FastAPI when the deliverable is the API itself: microservices, ML/LLM serving, streaming, and high-concurrency I/O-bound workloads, where async-native execution, Pydantic contracts, and auto-OpenAPI dominate. Senior nuance: the choice is rarely purely technical — team familiarity, existing infrastructure, and hiring matter; and hybrid architectures (Django for the admin-facing monolith, FastAPI for the inference microservice) are common and legitimate.

Q4. How does FastAPI generate its interactive documentation, and why is this more reliable than hand-written docs?

Answer. At startup FastAPI introspects every route: path, HTTP method, parameter annotations, Pydantic models, declared status codes, and security dependencies, and compiles them into an OpenAPI 3.x JSON document served at `/openapi.json`. Swagger UI (`/docs`) and ReDoc (`/redoc`) are just static frontends rendering that document. Reliability follows from provenance: the schema is *derived from* the code that executes, so it cannot drift the way a manually maintained wiki page does. Senior addition: the OpenAPI artifact is also a contract-testing and SDK-generation input in CI — documentation accuracy becomes a build guarantee, not a discipline problem.

Q5. A `sync def predict()` spends 200ms in-process; an `async def fetch()` awaits a 200ms upstream call. How does FastAPI schedule each, and what are the throughput implications?

Answer. FastAPI runs plain `def` handlers in a *threadpool* (via Starlette's `run_in_threadpool`), so the CPU-bound 200 ms model call occupies a thread but does not block the event loop; throughput is bounded by threadpool size and, for true CPU work, by the GIL and core count. The `async def` handler runs *on* the event loop; awaiting the upstream call suspends the coroutine, freeing the loop to serve other requests, so thousands of such requests can be in flight concurrently. The classic failure mode: writing `async def` but making a *blocking* 200 ms call inside it — this freezes the entire event loop and collapses throughput for every endpoint. Rule: `async` + `await` for I/O; plain `def` (or explicit offloading) for CPU.

Q6. What is OpenAPI, and how does it differ from Swagger UI?

Answer. OpenAPI is a vendor-neutral specification format — a JSON/YAML document describing an API's paths, schemas, parameters, and security. Swagger UI is one particular *viewer*: a web app that renders an OpenAPI document as interactive documentation. “Swagger” was the original name of the spec before it was donated and renamed OpenAPI (v3); the tooling brand kept the old name, which is why the terms are commonly conflated.

Q7. Why are Python type hints central to FastAPI rather than just stylistic?

Answer. In most Python code, type hints are passive metadata for linters. In FastAPI they are *executed*: the framework reads `item_id: int` at startup and from that one annotation derives request parsing, type coercion, validation with structured 422 errors, OpenAPI schema entries, and docs rendering. The annotation is the single source of truth for the request contract, eliminating the duplication (validation code + docs + serializer) that other stacks require.

Q8. What does “statelessness” mean in REST, and what does it buy you operationally?

Answer. Each request must carry everything the server needs to process it (credentials, identifiers, parameters); the server keeps no per-client session memory between requests. Operationally this is what makes horizontal scaling trivial: any worker can serve any request, so load balancers need no sticky sessions, instances can be added/removed/replaced freely, and a crashed node loses no client state. State that must persist moves to shared stores (database, Redis) — a theme that recurs in Chapters 5 and 10.

Q9. What is Uvicorn, and why does FastAPI need a separate server at all?

Answer. FastAPI is an ASGI *application* — a callable that consumes ASGI events. It cannot open sockets, parse raw HTTP, or manage the event loop; that is the job of an ASGI *server*. Uvicorn is the standard choice: it listens on the network, parses HTTP (with optional `uvloop/httptools` accelerators), translates each connection into ASGI events, and invokes the application. The separation mirrors WSGI’s app/server split (Flask/Gunicorn) and lets servers and frameworks evolve independently.

Q10. Your team’s hand-maintained API docs keep drifting from reality and integration partners complain. How would FastAPI change this workflow?

Answer. Move the contract into code: Pydantic models and typed endpoint signatures become the only place the API shape is defined. FastAPI then publishes `/openapi.json` that is correct by construction, with Swagger UI for exploration. In CI, snapshot the OpenAPI document and fail the build on uncommunicated breaking changes; optionally generate client SDKs from it so partners consume generated, always-current bindings. The process change is the point: documentation stops being a parallel artifact to maintain and becomes a build output.

Python Typing, Pydantic, and Data Validation

Every byte that enters or leaves your API passes through a validation boundary. This chapter builds deep fluency with that boundary: Python's typing system, Pydantic v2 models, request and response schemas, nested structures, field-level validation, and — critically for production — how to guarantee sensitive data never leaks out of a response.

2.1 Python Type Hints: The Substrate

Python's type annotations (PEP 484 and successors) attach type metadata to names: `def f(x: int) -> str`. The interpreter ignores them at runtime — they are stored, not enforced. Two kinds of tools give them power: *static* checkers (mypy, pyright) that analyze code before it runs, and *runtime* consumers that introspect annotations via `typing.get_type_hints()` and act on them. Pydantic is the most important runtime consumer in the Python ecosystem, and FastAPI is built on top of it.

The vocabulary you will use constantly: `Optional[str]` (i.e. `str | None`), container generics like `list[int]` and `dict[str, float]`, `Literal["a", "b"]` for enumerated values, `Annotated[T, metadata]` for attaching validation metadata to a type, and Union types for polymorphic payloads.

• CORE CONCEPT — Parse, don't validate

Pydantic's philosophy: instead of checking raw data and passing it along still-raw ("validate"), *convert* it into a typed object that makes illegal states unrepresentable ("parse"). After `User.model_validate(data)` succeeds, every consumer downstream can trust `user.age` is an `int` within bounds — no defensive re-checking. The boundary does the work once; the core stays clean.

2.2 BaseModel: Declaring Schemas

A Pydantic model is a class inheriting `BaseModel` whose annotated class attributes define fields:

```
1 from datetime import datetime
2 from typing import Literal
3 from pydantic import BaseModel, EmailStr, Field
```

```

4
5 class UserCreate(BaseModel):
6     """Inbound schema: what clients may send."""
7     email: EmailStr                                # required
8     password: str = Field(min_length=12)            # required + constraint
9     display_name: str | None = None                 # optional
10    role: Literal["viewer", "editor"] = "viewer"     # default + enum
11    tags: list[str] = Field(default_factory=list)     # mutable default
12
13 class UserOut(BaseModel):
14     """Outbound schema: what the API exposes. No password, ever."""
15     id: int
16     email: EmailStr
17     display_name: str | None
18     role: str
19     created_at: datetime

```

Listing 2.1: Request and response schemas with required, optional, and defaulted fields.

The required/optional rule is simple: *a field with no default is required; a field with a default is optional*. Note `default_factory=list` — never use a mutable literal as a default. Note also the deliberate asymmetry: `UserCreate` and `UserOut` are *different models*. Inbound and outbound schemas almost always diverge (passwords in, ids and timestamps out), and modeling them separately is the first habit of professional API design.

2.3 Request and Response Schemas in FastAPI

A Pydantic-typed parameter becomes the request body; `response_model` declares and *enforces* the response shape:

```

1 from fastapi import FastAPI, status
2
3 app = FastAPI()
4
5 @app.post(
6     "/users",
7     response_model=UserOut,
8     status_code=status.HTTP_201_CREATED,
9 )
10 def create_user(payload: UserCreate) -> UserOut:
11     user = user_service.create(payload) # returns ORM object
12     return user # FastAPI filters it through UserOut

```

Listing 2.2: The validation boundary in both directions.

Three things happen on the inbound side: the JSON body is parsed, coerced ("42" → 42 in lax mode), and validated; any failure short-circuits into a 422 response listing every invalid field with its location (body → password) and reason — your handler never runs with bad data. On the outbound side, `response_model` acts as a **whitelist**: whatever object you return, only fields declared on `UserOut` are serialized.




Figure 2.1: The two-way validation boundary: Pydantic guards what enters the handler; `response_model` guards what leaves it.

2.4 Nested Models and Field Validation

Real payloads are trees. Pydantic validates recursively — a model field typed as another model, or a list of models, is parsed all the way down:

```

1  from pydantic import BaseModel, Field, field_validator, model_validator
2
3  class LineItem(BaseModel):
4      sku: str = Field(pattern=r"^[A-Z]{3}-\d{4}$")
5      quantity: int = Field(gt=0, le=1_000)
6      unit_price: float = Field(gt=0)
7
8  class Address(BaseModel):
9      street: str
10     city: str
11     country: str = Field(min_length=2, max_length=2) # ISO code
12
13  class OrderCreate(BaseModel):
14     customer_email: str
15     items: list[LineItem] = Field(min_length=1)
16     shipping: Address
17     billing: Address | None = None
18
19     @field_validator("customer_email")
20     @classmethod
21     def normalize_email(cls, v: str) -> str:
22         return v.strip().lower()
23
24     @model_validator(mode="after")
25     def billing_defaults_to_shipping(self) -> "OrderCreate":
26         if self.billing is None:
27             self.billing = self.shipping
28         return self
  
```

Listing 2.3: Nested models with field and model-level validators (Pydantic v2).

`Field()` carries declarative constraints (`gt`, `le`, `min_length`, `pattern`); `@field_validator` adds custom per-field logic; `@model_validator(mode="after")` enforces cross-field invariants on the constructed object. All constraints flow into the OpenAPI schema automatically, so the docs show them to consumers.

2.5 Pydantic v2 Essentials

Pydantic v2 (2023) rewrote the validation core in Rust (`pydantic-core`), delivering 5–50× faster validation — material at high request rates. The API changed with it; this book uses v2 idioms throughout:

- `model_validate()` / `model_dump()` / `model_dump_json()` replace v1's `parse_obj()` / `dict()` / `json()`.
- `@field_validator` and `@model_validator` replace `@validator` and `@root_validator`.
- Configuration moves to `model_config = ConfigDict(...)` — e.g. `from_attributes=True` to read ORM objects, and `extra="forbid"` to reject unknown fields.
- **Strict vs lax mode:** lax (default) coerces compatible types ("5" → 5); strict mode rejects them. Choose per-field via `Field(strict=True)` when coercion would mask client bugs.

2.6 Preventing Sensitive Data Leaks

The most expensive validation bug is not rejecting bad input — it is *emitting* data you never meant to expose. Returning an ORM user object from an endpoint without a response model serializes *every* column: password hash, internal flags, soft-delete markers.

Never rely on the caller of `model_dump()` to remember `exclude={...}`. Exclusion at call sites is a convention; one forgotten call site is a breach. Make leak-prevention *structural*: dedicated response models that simply do not declare sensitive fields.

```

1 from pydantic import BaseModel, ConfigDict, SecretStr
2
3 class UserDB(BaseModel):
4     model_config = ConfigDict(from_attributes=True)
5     id: int
6     email: str
7     hashed_password: str
8     is_superuser: bool
9     stripe_customer_id: str
10
11 class UserPublic(BaseModel):
12     """Defense 1: whitelist-by-design. Only these fields can ever leave."""
13     model_config = ConfigDict(from_attributes=True)
14     id: int
15     email: str
16
17 class LoginRequest(BaseModel):
18     email: str
19     # Defense 2: SecretStr masks the value in logs/repr -> '*****'
20     password: SecretStr
21
22 @app.get("/users/{user_id}", response_model=UserPublic)
23 def get_user(user_id: int):
24     return repo.get(user_id) # full ORM object in, two fields out

```

Listing 2.4: Layered defenses against data leakage.**PRODUCTION CASE STUDY — The Rails-era mass-assignment lesson, replayed in APIs**

In 2012 a researcher exploited GitHub’s Ruby on Rails mass-assignment defaults to add his own public key to the Rails organization — the inbound payload was bound directly to the model, so posting an unexpected field silently set it. The same class of bug appears today in FastAPI services that reuse one model for input, storage, and output: a client POSTs `{"role": "admin"}`, the shared model has a `role` field, and privilege escalation follows. Production teams therefore enforce the *three-model rule* — `XCreate` (inbound, `extra="forbid"`), `XDB` (storage), `XPublic` (outbound) — and add a CI check that no endpoint lacks a `response_model`. The fix is architectural, not a patch: separate schemas make both mass assignment and field leakage impossible by construction.

COST MODEL & METRICS — The cost of weak validation

Let N = requests/day, p = fraction of malformed payloads reaching business logic, and $C_{incident}$ = average cost of a data-quality incident (engineer-hours, reprocessing, customer impact). Expected monthly cost:

$$C_{monthly} = 30 \cdot N \cdot p \cdot q \cdot C_{incident}$$

where q is the probability a malformed payload that slips through causes a real incident. With $N = 10^6$, $p = 10^{-4}$, $q = 10^{-2}$, and $C_{incident} = \$2,000$: \$60,000/month of expected loss — against a validation layer that costs microseconds per request. Boundary validation is one of the highest-ROI lines of code you will ever write. Measure it: track `422_rate` (clients sending garbage) and `5xx_after_validation` (garbage that got through) as first-class metrics.

SYSTEM DESIGN SCENARIO — Schema design for a multi-tenant document-ingestion API

Requirements: customers upload documents (title, body, metadata dict, up to 200 tags) for indexing; tenants must never see each other’s fields; the ML team adds new optional metadata monthly without breaking clients.

Reference approach: (1) `DocumentCreate` with `extra="forbid"` and hard caps (body max length, tags max_length=200) — caps are also cost controls on the downstream embedding pipeline. (2) `tenant_id` comes from the auth token (Chapter 6), *never* from the body — any client-supplied tenant field is rejected by `extra="forbid"`. (3) `DocumentPublic` response model omits internal scoring fields. (4) Additive evolution: new fields are optional-with-default, so old clients keep working; renames/removals go through `/api/v2` (Chapter 3). (5) Version the schemas in code review: a CI snapshot of `/openapi.json` flags any accidental breaking change.

2.7 Interview Questions

Q1. Python type hints are not enforced by the interpreter. How does Pydantic make them enforceable, and what is the cost?

Answer. Pydantic introspects a model's annotations at *class-definition time* and compiles a per-model validation schema (in v2, executed by the Rust `pydantic-core` engine). At instantiation, raw data is parsed against that compiled schema — coercing, validating constraints, and raising structured `ValidationError`s. So enforcement happens at explicit parsing boundaries, not pervasively. The cost is per-instantiation CPU: microseconds for typical payloads in v2, but it scales with payload size and nesting depth, so megabyte-scale deeply nested bodies deserve benchmarks and size caps. Senior nuance: enforcement is *boundary validation* — internal function calls remain unchecked, which is exactly the “parse, don't validate” design: validate once at the edge, trust types inside.

Q2. Why should request, storage, and response schemas be separate models even though it triples the class count?

Answer. Because the three contracts genuinely differ and coupling them creates two vulnerability classes. Inbound: a shared model accepts fields the client should not control (`role`, `is_admin`) — mass assignment. Outbound: a shared model emits fields the client should not see (`hashed_password`, internal IDs) — data leakage. Separate models make both impossible structurally rather than by call-site discipline. They also decouple evolution: you can add a stored column without changing the public contract, or accept a new input field without persisting it. Boilerplate is manageable with inheritance (a shared `UserBase`) and is trivially cheaper than one leaked credential field.

Q3. Explain Pydantic v2's lax vs strict modes. When is coercion dangerous?

Answer. Lax mode (default) coerces compatible inputs: `"42" → 42`, `1 → True`, numeric strings to floats. Strict mode rejects any type mismatch. Coercion is convenient for query/path parameters (which arrive as strings by nature) but dangerous where type distinctions carry meaning: financial amounts (`"100"` vs `100` may signal a client serialization bug worth surfacing, and float coercion of monetary strings invites rounding errors — use `Decimal`), booleans (lax accepts `"yes"/1`, masking malformed clients), and IDs where `int/str` confusion can cross-reference the wrong entity. Pragmatic policy: lax at the HTTP edge, `Field(strict=True)` on high-stakes fields, strict mode for internal service-to-service contracts where both sides are typed.

Q4. An endpoint returns the SQLAlchemy User object directly and the response contains hashed_password. Walk through the layers of defense you would add.

Answer. Immediate fix: declare `response_model=UserPublic` where `UserPublic` whitelists only public fields (`from_attributes=True` to read the ORM object). Structural defenses: (1) three-model separation so no outbound path can reference sensitive fields; (2) `SecretStr` for credentials so even logging/repr masks them; (3) a CI lint that fails any route lacking an explicit `response_model`; (4) a contract test asserting the serialized response keys equal the public schema exactly; (5) optionally, a response middleware in staging that scans for known-sensitive key names as a canary. The senior point: treat the incident as a process gap (nothing forced a whitelist), not a one-line bug.

Q5. How do nested model validation errors surface to the client, and why is the error structure important for API consumers?

Answer. Pydantic collects *all* failures (not just the first) with a `loc` path into the structure — e.g. `["body", "items", 2, "quantity"]` — plus a machine-readable `type` and human message; FastAPI wraps this list in a 422 response. The structure matters because consumers can programmatically map each error to a UI field (form highlighting), retry intelligently, and localize messages from the `type` code rather than parsing English strings. Collecting all errors at once also halves integration round-trips versus fail-fast validators. Senior addition: standardize on this error envelope across your services so client error-handling is written once; and avoid leaking internals by overriding the default handler if your field names are sensitive.

Q6. What is the difference between a required field, an optional field, and a nullable field in Pydantic?

Answer. Required vs optional is about *presence*: a field with no default must appear in the input; one with a default may be omitted. Nullable is about *value*: `str | None` permits an explicit JSON `null`. They combine independently — `name: str | None` (no default) is *required but nullable*: the key must be present, though its value may be null; `name: str = "x"` is optional but non-nullable. Confusing the two produces subtle contract bugs, especially for PATCH endpoints where “omitted” means “leave unchanged” but “null” means “clear the value.”

Q7. What does `Field(default_factory=list)` solve, and why not write `tags: list[str] = []`?

Answer. In ordinary Python, a mutable default is created once and shared by all calls — the classic shared-mutable-default bug. Pydantic actually guards against this internally (it deep-copies defaults), but `default_factory` is the explicit, idiomatic form: a zero-argument callable invoked per instantiation, guaranteeing a fresh object. It is also the only correct option for

dynamic defaults like `default_factory=datetime.utcnow` or `uuid4`, where the value must be computed at creation time rather than at class-definition time.

Q8. How do Pydantic `Field` constraints interact with the OpenAPI docs?

Answer. Every declarative constraint — `gt`, `ge`, `le`, `min_length`, `max_length`, `pattern`, plus `title`/`description`/`examples` — is exported into the JSON Schema embedded in `/openapi.json`, and Swagger UI renders them on each field. Consumers see “quantity: integer, > 0, ≤ 1000” without any separate documentation. Custom `@field_validator` logic, by contrast, is opaque to the schema — so prefer declarative constraints where possible and document validator behavior in the field’s description.

Q9. When would you use `@model_validator` instead of `@field_validator`?

Answer. `@field_validator` sees one field at a time, so it suits normalization and single-field rules (trim, lowercase, format checks). `@model_validator(mode="after")` runs on the fully constructed model and is required for *cross-field invariants*: `end_date > start_date`, “exactly one of `file_url` or `inline_content` must be set,” or defaulting one field from another (billing := shipping). `mode="before"` variants run on the raw input dict and are useful for reshaping legacy payloads before field validation.

Q10. Your API’s p99 latency budget is 50ms and payloads are 2MB nested JSON. How do you keep validation from blowing the budget?

Answer. First measure: profile `model_validate` on representative payloads — Pydantic v2’s Rust core handles most 2MB documents in single-digit milliseconds, so the problem may not exist. If it does: (1) impose size caps at the edge (reject oversized bodies before parsing); (2) flatten or cap pathological nesting/array lengths via `max_length` constraints; (3) skip double validation — don’t re-validate trusted internal data (`model_construct()` builds without validation when provenance is safe); (4) avoid heavyweight custom validators doing I/O (move those to the service layer); (5) for extreme hot paths, validate lazily only the fields the endpoint touches. Keep the boundary — never remove validation on untrusted input to save milliseconds; restructure it instead.

Building Core API Endpoints

This chapter turns schema knowledge into working APIs: the full set of HTTP verbs, path and query parameters, status codes, CRUD design, pagination, filtering, sorting, and versioning. The running example — a task-management API — is deliberately boring, because the patterns are what transfer to every service you will ever build, including the AI services later in this book.

3.1 Verbs, Semantics, and Idempotency

Each HTTP method carries semantics that clients, proxies, and caches rely on. GET is *safe* (no side effects) and cacheable. PUT and DELETE are *idempotent*: applying them twice yields the same state as once — which is what makes client retries safe. POST is neither safe nor idempotent; a retried POST may create a duplicate, which is why payment APIs add idempotency keys. PATCH applies a partial update; PUT replaces the whole resource.

• CORE CONCEPT — Idempotency is a retry contract

Distributed systems retry. Networks drop responses after the server has already acted. Designing PUT/DELETE to be idempotent — and giving POST an idempotency-key escape hatch — is not pedantry; it is what makes “client timed out, retried, and now there are two orders” impossible. Interviewers probe this constantly.

3.2 Path Parameters, Query Parameters, and Bodies

The three input channels divide cleanly by purpose: **path** parameters identify *which resource* (/tasks/42); **query** parameters modify *how* the collection is read (filters, sort, pagination); the **body** carries resource *content* on writes.

```
1 from enum import Enum
2 from typing import Annotated
3 from fastapi import APIRouter, HTTPException, Path, Query, status
4 from pydantic import BaseModel, Field
5
6 router = APIRouter(prefix="/api/v1/tasks", tags=["tasks"])
7
8 class TaskStatus(str, Enum):
9     todo = "todo"
```

```

10     in_progress = "in_progress"
11     done = "done"
12
13     class TaskCreate(BaseModel):
14         title: str = Field(min_length=1, max_length=200)
15         description: str = ""
16         status: TaskStatus = TaskStatus.todo
17
18     class TaskUpdate(BaseModel):                # PATCH: everything optional
19         title: str | None = Field(default=None, min_length=1, max_length=200)
20         description: str | None = None
21         status: TaskStatus | None = None
22
23     class TaskOut(BaseModel):
24         id: int
25         title: str
26         description: str
27         status: TaskStatus
28
29     class Page(BaseModel):
30         items: list[TaskOut]
31         total: int
32         limit: int
33         offset: int
34
35     @router.post("", response_model=TaskOut,
36                 status_code=status.HTTP_201_CREATED,
37                 summary="Create a task")
38     def create_task(payload: TaskCreate):
39         return repo.create(payload)
40
41     @router.get("", response_model=Page, summary="List tasks")
42     def list_tasks(
43         status_: Annotated[TaskStatus | None, Query(alias="status")] = None,
44         q: Annotated[str | None, Query(max_length=100,
45                                     description="Full-text search in title")] = None,
46         sort: Annotated[str, Query(pattern=r"^-?(id|title|status)$")] = "id",
47         limit: Annotated[int, Query(ge=1, le=100)] = 20,
48         offset: Annotated[int, Query(ge=0)] = 0,
49     ):
50         items, total = repo.list(status_, q, sort, limit, offset)
51         return Page(items=items, total=total, limit=limit, offset=offset)
52
53     @router.get("/{task_id}", response_model=TaskOut)
54     def get_task(task_id: Annotated[int, Path(ge=1)]):
55         task = repo.get(task_id)
56         if task is None:
57             raise HTTPException(status_code=404, detail="Task not found")
58         return task
59
60     @router.put("/{task_id}", response_model=TaskOut,
61                 summary="Replace a task entirely")
62     def replace_task(task_id: int, payload: TaskCreate):
63         return repo.replace(task_id, payload)
64
65     @router.patch("/{task_id}", response_model=TaskOut,
66                  summary="Partially update a task")
67     def update_task(task_id: int, payload: TaskUpdate):

```

```

68     # exclude_unset: only fields the client actually sent
69     changes = payload.model_dump(exclude_unset=True)
70     return repo.update(task_id, changes)
71
72 @router.delete("/{task_id}", status_code=status.HTTP_204_NO_CONTENT)
73 def delete_task(task_id: int):
74     repo.delete(task_id)

```

Listing 3.1: A complete, production-shaped CRUD router for tasks.

Several production details deserve attention. The PATCH handler uses `model_dump(exclude_unset=True)` — the canonical way to distinguish “client omitted the field” from “client sent null.” Query constraints (`le=100` on `limit`) are cost controls, not decoration: an unbounded `limit` is an invitation to dump your table. The `sort` pattern whitelist prevents callers from injecting arbitrary column names into your SQL layer.

3.3 Status Codes That Mean What They Say

Use 201 Created for successful creation (with a `Location` header when practical), 204 No Content for deletes, 400 for semantically invalid requests, 401 for missing/invalid credentials, 403 for valid credentials lacking permission, 404 for absent resources, 409 Conflict for state collisions (duplicate unique key, concurrent edit), 422 for schema validation failures (FastAPI’s default), and 429 for rate limiting. Never return 200 with an `{"error": ...}` body — it breaks caches, monitoring, retry logic, and every generic HTTP client ever written.

3.4 Pagination, Filtering, and Sorting

Offset pagination (`limit/offset`) is simple and supports random access, but degrades on deep pages (the database must scan and discard `offset` rows) and can skip/duplicate items when rows are inserted mid-scan. **Cursor (keyset) pagination** encodes the last-seen sort key (e.g. `?after=2024-06-01,9182`) and uses an indexed `WHERE (created_at, id) > (...)` filter — constant cost at any depth and stable under writes. Rule of thumb: offset for small admin UIs, cursor for anything public, infinite-scroll, or large.




Figure 3.1: URL-based API versioning: both versions route into one service layer; v1 survives through a thin adapter, so business logic is never duplicated.

3.5 API Versioning with /api/v1

Breaking changes — removing a field, renaming, changing types or semantics — require a new version; additive optional fields do not. URL-path versioning (`/api/v1/...`) is the most common because it is visible, cacheable, and trivially routable. In FastAPI, mount one `APIRouter` per version (Figure 3.1); keep business logic in a shared service layer and let the old version live as a thin adapter. Publish a deprecation window (e.g. 6–12 months), emit `Deprecation` and `Sunset` headers from v1, and track v1 traffic so you know when it is safe to remove.

PRODUCTION CASE STUDY — Stripe — the gold standard of versioned CRUD design

Stripe’s API is the canonical demonstration that disciplined CRUD mechanics compound into a business asset. Every list endpoint is cursor-paginated (`starting_after/ending_before`) because offset pagination over billions of payment rows is both slow and unstable. Every POST accepts an `Idempotency-Key` header, making client retries safe in the one domain where duplicates mean charging a card twice. And versioning is date-stamped per account with server-side compatibility shims, so integrations written a decade ago still run. The transferable lesson: pagination strategy, idempotency, and versioning are not polish you add later — they are load-bearing decisions that determine whether your API can be trusted with money. Design them in Chapter-3 of your project, not year three.

COST MODEL & METRICS — The deep-pagination cost curve

Offset pagination cost grows linearly with depth: fetching page k of size n forces the database to process $\Theta(k \cdot n)$ rows for n returned. Total rows processed for a client scanning P pages:

$$W_{offset} = \sum_{k=1}^P k n = \frac{P(P+1)}{2} n \quad \text{vs} \quad W_{cursor} = P n.$$

A crawler reading 1,000 pages of 100 rows costs $\sim 50\text{M}$ row-visits with offset but 100K with cursors — a $500\times$ difference in database work for identical output. Metrics to watch: p95 latency *by page depth*, and DB rows-examined/rows-returned ratio (a ratio $\gg 1$ on list endpoints usually means offset pagination or missing indexes).

SYSTEM DESIGN SCENARIO — Design the read API for a product catalog (10M SKUs)

Requirements: faceted filtering (category, brand, price range), sorting by price/popularity, 2,000 req/s reads, p95 < 100 ms, mobile clients on flaky networks.

Reference approach: (1) GET `/api/v1/products` with whitelisted, validated query params; every filterable column backed by an index (composite indexes matching the hot filter+sort combinations). (2) Cursor pagination keyed on (`sort_field`, `id`) for stability; expose `next_cursor` opaque token, never raw offsets. (3) Cap limit at 50; mobile clients page incrementally. (4) ETag headers on detail endpoints so clients revalidate cheaply. (5) Read replicas or a search engine (OpenSearch) once facet combinations exceed what B-tree indexes serve; the API contract stays identical — only the repository implementation changes (Chapter 4’s layering pays off

here). (6) 429 + Retry-After protects against crawler storms.

3.6 Interview Questions

Q1. Explain idempotency for each HTTP verb, and how you would make POST safe to retry.

Answer. GET/HEAD are safe and idempotent (no state change). PUT is idempotent: replacing a resource with the same representation twice leaves identical state. DELETE is idempotent in *effect* (resource ends absent), though the second call may return 404. PATCH is not guaranteed idempotent (e.g. “append to list” patches). POST is not idempotent — retries can duplicate. The fix is an **idempotency key**: the client sends a unique header per logical operation; the server atomically records the key before processing and returns the stored response on replays. Senior details: the key store needs a TTL and must be checked within the same transaction/lock as the side effect, otherwise two concurrent retries race; and the stored response must include the original status code.

Q2. Offset vs cursor pagination: trade-offs, failure modes, and when each is correct.

Answer. Offset (LIMIT n OFFSET m) gives random page access and trivial implementation, but the database walks past m rows each query (linear degradation), and concurrent inserts/deletes shift rows between pages, causing duplicates or gaps. Cursor/keyset encodes the last-seen sort key and queries WHERE (key, id) > (cursor) on an index — $O(\log N)$ at any depth and stable under writes — but loses random access and requires a deterministic, indexed, *unique* sort (hence the id tiebreaker). Offset is fine for small, internal, jump-to-page UIs; cursor is correct for public APIs, feeds, and any large/hot table. Senior addition: make cursors opaque (base64 of key material) so clients cannot construct or tamper with them, and version the cursor format.

Q3. A client PATCHes {"description": null}. How do you distinguish “clear this field” from “field not provided,” and how does Pydantic support it?

Answer. Three client intents exist: field omitted (leave unchanged), field null (clear it), field valued (set it). JSON cannot distinguish the first two by value alone — you need *presence* information. Pydantic tracks which fields were actually supplied; `model_dump(exclude_unset=True)` returns only provided fields, so an omitted `description` never appears, while an explicit null appears as `None`. The update layer then applies exactly those keys. Senior nuances: every PATCH-model field must be optional with no semantic default (a default would be indistinguishable from client input); and if “clear” is disallowed for some field, validate `None` explicitly rather than silently ignoring it.

Q4. Walk through correct status-code choices for: duplicate signup email, expired token, valid token but wrong role, malformed JSON, schema-valid but business-rule-violating payload, and a deleted resource.

Answer. Duplicate email: 409 `Conflict` (state collision with an existing resource) — though some teams return 200-with-generic-message on signup specifically to avoid account enumeration; flagging that trade-off is a senior signal. Expired token: 401 with `WWW-Authenticate` — authentication failed. Wrong role: 403 — authenticated but not authorized. Malformed JSON: 400 (FastAPI returns 422 from its body parser; either is defensible, consistency matters). Schema-valid but business-invalid (e.g. withdrawing more than balance): 422 or 409 depending on whether it is an input problem or a state problem. Deleted resource: 404, or 410 `Gone` if you want to signal permanence. The meta-answer interviewers want: codes are a contract consumed by machines — retry logic treats 4xx as “don’t retry” and 5xx as “retry,” so miscoding errors breaks clients you have never met.

Q5. How would you evolve an API from v1 to v2 without breaking existing consumers? Describe the full lifecycle.

Answer. (1) Classify the change: additive optional fields need no version; renames, removals, type or semantic changes do. (2) Mount `/api/v2` as a new router; refactor shared business logic into the service layer so v1 becomes a thin adapter over the same core — never fork the logic. (3) Announce with a deprecation policy; emit `Deprecation` and `Sunset` headers on v1 responses and document the migration. (4) Instrument per-version traffic and identify remaining v1 consumers; contact the long tail directly. (5) Brown-out v1 (scheduled short failures) before final removal to flush silent dependents. (6) Remove v1; keep the OpenAPI history archived. Senior signals: contract tests pinning v1 behavior while v2 evolves, and resisting version proliferation — batch breaking changes rather than minting v3, v4 casually.

Q6. Why must list-endpoint `limit` parameters have a server-enforced maximum?

Answer. An unbounded limit lets any caller (or bug) request the entire table: memory spikes on server and database, slow serialization blocking the worker, and a trivially exploitable denial-of-service vector. A cap (`Query(limit=100)`) bounds the worst-case cost of a single request, which is the foundation of capacity planning — you cannot size a system whose per-request cost is unbounded. It also surfaces in OpenAPI so clients learn the contract upfront rather than via production failures.

Q7. What belongs in a path parameter vs a query parameter, and why does the distinction matter?

Answer. Path parameters identify the resource — they are part of its name (`/tasks/42`); a different value means a different resource, and a missing one is structurally impossible. Query parameters parameterize the *operation* on a resource or collection: filters, sort, pagination, projection — optional by nature, with defaults. The distinction matters for caching (URLs are cache keys), for 404 semantics (`/tasks/999` not existing is a 404; `?status=bogus` is a 422), and for OpenAPI clarity. A smell worth naming: filters in the path (`/tasks/done/recent`) explode the route space and break composability.

Q8. What is the practical difference between PUT and PATCH for an update endpoint?

Answer. PUT replaces the entire resource: the request body is the complete new representation, so omitted fields are reset to defaults — and the request model equals the create model. PATCH applies a partial change: only provided fields are touched, requiring an all-optional model and `exclude_unset` handling. PUT is idempotent by definition; PATCH usually is in practice (set-style patches) but is not guaranteed. Offering PUT when clients routinely send partial bodies is a classic data-loss bug: the omitted fields get silently wiped.

Q9. How do tags, summaries, and descriptions on endpoints affect more than just aesthetics?

Answer. They structure the OpenAPI document, which downstream tooling consumes: Swagger UI groups operations by tag (navigability for human integrators), SDK generators turn tags into client namespaces/classes and summaries into method docstrings, and API gateways/portals use them for catalog organization. Inconsistent tagging therefore produces badly shaped generated SDKs — a real integration cost, not a cosmetic one. Teams typically standardize one tag per resource/router and enforce summary presence in review or lint.

Q10. Your list endpoint is p95-slow only for a handful of consumers. Investigation shows they paginate to page 4,000. What are your options?

Answer. Confirm the mechanism: offset pagination scanning `offset × limit` rows. Options, in preference order: (1) migrate to cursor pagination — constant-cost depth; keep offset params working for shallow pages during transition. (2) If consumers are doing bulk export, give them a purpose-built path: an async export job (Chapter 8) producing a file, or a changes-since feed (`?updated_after=`) — bulk-via-pagination is usually a requirements smell. (3) Cap maximum page depth and document it. (4) Add the covering index that makes the keyset query cheap. The senior framing: deep pagination is almost always a misused workload; fix the workload, not just the query.

Dependency Injection and Application Structure

The difference between a demo and a maintainable service is structure. This chapter covers FastAPI's dependency injection system — `Depends()` — and the layered architecture it enables: routers, services, repositories, and the project layout that keeps a growing codebase testable. These patterns are prerequisites for every remaining chapter.

4.1 The Dependency Injection Concept

Dependency injection (DI) inverts the question “where does my handler get its database session, current user, or configuration?” Instead of the handler constructing or importing these collaborators (coupling it to concrete implementations), it *declares* what it needs, and the framework supplies them. Three benefits follow: **decoupling** (handlers depend on interfaces, not constructions), **reuse** (one `get_db` serves every endpoint), and — most important — **substitutability**: tests can replace any dependency without patching internals (Chapter 7 builds entirely on this).

In FastAPI, a dependency is just a callable. `Depends(get_db)` tells the framework: before calling this handler, call `get_db` (resolving its declared dependencies recursively), and pass the result in.

```

1  from typing import Annotated
2  from fastapi import Depends, FastAPI, Header, HTTPException
3
4  app = FastAPI()
5
6  # --- 1. Resource with cleanup: yield-style dependency -----
7  def get_db():
8      db = SessionLocal()
9      try:
10         yield db          # handler runs while we are suspended here
11     finally:
12         db.close()        # always runs, even on exceptions
13
14  DB = Annotated[Session, Depends(get_db)]
15
16  # --- 2. Chained dependencies: auth built on top of the db -----
17  def get_current_user(db: DB, authorization: str = Header(...)) -> User:
18      user = decode_and_load_user(db, authorization)
19      if user is None:

```

```

20         raise HTTPException(status_code=401)
21     return user
22
23     CurrentUser = Annotated[User, Depends(get_current_user)]
24
25     def require_admin(user: CurrentUser) -> User:
26         if user.role != "admin":
27             raise HTTPException(status_code=403)
28         return user
29
30     # --- 3. Handlers declare; framework supplies -----
31     @app.get("/me")
32     def read_me(user: CurrentUser):
33         return {"email": user.email}
34
35     @app.delete("/users/{user_id}", dependencies=[Depends(require_admin)])
36     def delete_user(user_id: int, db: DB):
37         ...

```

Listing 4.1: Core dependency patterns: yield-cleanup, chaining, and caching.

Three mechanics matter. **Yield dependencies** wrap the request in a setup/teardown pair — the idiomatic home for sessions, connections, and locks. **Chaining:** `get_current_user` itself depends on `get_db`; FastAPI resolves the whole graph. **Per-request caching:** if five dependencies in one request all require `get_db`, it executes *once* and the result is shared (disable with `Depends(get_db, use_cache=False)` when you genuinely need separate instances).

4.1.1 Router-Level and App-Level Dependencies

Dependencies attach at three scopes. *Parameter-level* supplies a value to one handler. *Router-level* (`APIRouter(dependencies=[Depends(verify_api_key)])`) runs for every route in the router — ideal for “everything under `/admin` requires admin.” *App-level* (`FastAPI(dependencies=[...])`) runs for every request — a lightweight alternative to middleware when the check needs the DI system (middleware cannot use `Depends`).

4.2 Routers and Project Structure

`APIRouter` is the unit of modularity: each resource gets its own router file, and `main.py` shrinks to composition. The layout below scales from three endpoints to three hundred:

```

1  app/
2  main.py                # create_app(), router mounting, lifespan
3  core/
4  config.py              # Settings (pydantic-settings), env-driven
5  security.py            # hashing, JWT helpers
6  api/
7  deps.py                # shared dependencies (DB, CurrentUser, ...)
8  v1/
9  tasks.py               # APIRouter for /api/v1/tasks
10 users.py               # APIRouter for /api/v1/users
11 services/
12 task_service.py        # business logic, transactions

```

```

13 repositories/
14     task_repo.py      # all SQL/ORM access for tasks
15 models/              # SQLAlchemy models
16 schemas/             # Pydantic models (Create/Update/Out)
17 tests/

```

Listing 4.2: A production FastAPI project layout.

4.3 Clean Architecture: Router, Service, Repository

The layering rule: **routers** translate HTTP ↔ domain (parse input, call service, map exceptions to status codes — no business logic); **services** hold business rules and orchestrate transactions (no HTTP, no raw SQL); **repositories** encapsulate all data access (no business rules). Dependencies point inward only: router → service → repository. The payoff is substitutability per layer: swap Postgres for a fake repository in service tests; exercise HTTP concerns without a database; move a service to a worker queue without touching SQL.




Figure 4.1: Layered FastAPI architecture. Each layer is independently replaceable; `Depends()` is the wiring between them.

```

1  # repositories/task_repo.py
2  class TaskRepository:
3      def __init__(self, db: Session):
4          self.db = db
5      def get(self, task_id: int) -> Task | None:
6          return self.db.get(Task, task_id)
7      def add(self, task: Task) -> Task:
8          self.db.add(task); self.db.flush(); return task
9
10 # services/task_service.py
11 class TaskService:
12     def __init__(self, repo: TaskRepository):
13         self.repo = repo
14     def complete_task(self, task_id: int, user: User) -> Task:
15         task = self.repo.get(task_id)
16         if task is None:
17             raise TaskNotFound(task_id)
18         if task.owner_id != user.id:
19             raise NotTaskOwner(task_id)           # domain exception
20         task.status = "done"
21         return task
22
23 # api/deps.py -- composition root

```

```

24 def get_task_repo(db: DB) -> TaskRepository:
25     return TaskRepository(db)
26
27 def get_task_service(
28     repo: Annotated[TaskRepository, Depends(get_task_repo)],
29 ) -> TaskService:
30     return TaskService(repo)
31
32 # api/v1/tasks.py -- thin router
33 @router.post("/{task_id}/complete", response_model=TaskOut)
34 def complete(
35     task_id: int,
36     user: CurrentUser,
37     svc: Annotated[TaskService, Depends(get_task_service)],
38 ):
39     try:
40         return svc.complete_task(task_id, user)
41     except TaskNotFound:
42         raise HTTPException(404, "Task not found")
43     except NotTaskOwner:
44         raise HTTPException(403, "Not your task")

```

Listing 4.3: The three layers wired together with class-based dependencies.

Note the service raises *domain* exceptions, not `HTTPException` — the service does not know HTTP exists, which is precisely what lets you reuse it from a CLI, a queue worker, or a gRPC front later. (Registering app-wide exception handlers for domain exceptions removes even the try/except boilerplate.)

PRODUCTION CASE STUDY — Refactoring a 40-endpoint monolith file at a fintech startup

A common (and commonly fatal) trajectory: a fintech team shipped its MVP with all endpoints, SQL, and business rules in one 3,000-line `main.py`. Symptoms at month nine: every PR conflicted with every other; tests required a live database and took 25 minutes; an engineer fixing a webhook accidentally changed loan-approval logic in a shared helper. The remediation followed exactly this chapter: (1) split routes into per-resource `APIRouters` — mechanical, zero-risk; (2) extract repositories, moving every raw query behind methods — enabling an SQL review that found three N+1 patterns and one injection risk; (3) extract services with domain exceptions; (4) introduce `deps.py` as the single composition root. Results they measured: test suite to 90 seconds (fake repositories), median PR review time halved, and the loan-logic module gained an owner and a test gate. The lesson is the cost curve: this refactor took two engineers three weeks at month nine; the same structure costs near-zero at week one.

COST MODEL & METRICS — Quantifying coupling: the change-amplification metric

Track *change amplification* A = average number of files touched per small feature PR, and test cost T = wall-clock CI minutes per run. In a layered codebase, adding one field end-to-end touches a bounded set (schema, model, repo, service, router $\Rightarrow A \approx 5$, independent of codebase size); in an unlayered one, A grows with the number of call sites, roughly $A \propto N_{endpoints}$.

Engineering cost of a feature scales as

$$C_{feature} \approx A \cdot (c_{edit} + c_{review}) + n_{runs} \cdot T,$$

so keeping A flat and T low (fast unit layers, few full-stack tests) is the economic argument for this chapter. Teams that measure A catch architectural erosion quarters before it becomes a rewrite.

SYSTEM DESIGN SCENARIO — Structure a multi-team FastAPI platform (12 engineers, 3 domains)

Requirements: billing, catalog, and notifications domains; teams deploy independently in the future but start as one service; shared auth; CI under 10 minutes.

Reference approach: (1) A modular monolith: one FastAPI app, `app/domains/{billing, catalog, notifications}/` each containing its own routers, services, repositories, schemas — domain folders may not import from each other; cross-domain calls go through explicit service interfaces defined in a shared `contracts/` package (enforced by an import-linter rule in CI). (2) Shared kernel: `core/` (config, security, DB engine) and `api/deps.py`. (3) Each domain mounts under `/api/v1/{domain}` via one top-level include per team. (4) Tests split per domain; CI runs changed domains' suites plus a thin cross-cutting smoke suite. (5) When a domain needs independent deployment, its folder is the extraction boundary: lift it into a new service, replace in-process contract calls with HTTP — the DI seams make this a move, not a rewrite.

4.4 Interview Questions

Q1. Explain how FastAPI resolves `Depends()` for a request, including chaining and caching.

Answer. At startup FastAPI builds a dependency graph per route by inspecting signatures recursively: a handler depends on `get_current_user`, which depends on `get_db`, etc. Per request, the graph is resolved depth-first: each dependency's own parameters (path/query/header/body or sub-dependencies) are supplied, it is called (sync dependencies in the threadpool, async on the loop), and its result is memoized in a per-request cache keyed by the callable — so multiple appearances of `get_db` share one session. Yield dependencies are entered like context managers and their teardown runs after the response, in reverse order. `use_cache=False` opts out of memoization. Senior addition: the graph is computed once at startup, so resolution overhead per request is small; and security dependencies participate in OpenAPI, marking routes as authenticated.

Q2. Why do yield-based dependencies matter for database session management? What happens on an unhandled exception?

Answer. A yield dependency brackets the request: code before `yield` acquires the resource, the handler runs while the generator is suspended, and the `finally` block releases it — guaranteeing the session closes (returning the connection to the pool) on success *and* failure. On an unhandled handler exception, FastAPI throws it back into the generator at the yield point, so `except/finally` clauses run — letting you roll back the transaction before closing. Without this, error paths leak connections; the pool exhausts under sustained errors and the whole service locks up — a classic production outage. Senior nuance: commit-on-success/rollback-on-error can live in the dependency itself, giving every endpoint transactional semantics for free.

Q3. Where should business logic live, and how do you keep HTTP concerns out of it? What goes wrong otherwise?

Answer. In the service layer: pure-Python classes/functions that take typed inputs, enforce invariants, orchestrate repositories and transactions, and raise *domain* exceptions. Routers stay thin translators: extract validated input, call the service, map domain exceptions to status codes. The enforcement mechanism is the exception boundary — the moment a service raises `HTTPException` or reads a `Request`, it is welded to HTTP and cannot be reused from a Celery worker, cron job, or test without faking a web stack. Other failure modes of logic-in-routers: untestable without `TestClient` round-trips, copy-paste divergence when a second consumer needs the same rule, and transaction scope smeared across HTTP plumbing. Senior signal: name the composition root — DI wiring lives in one place (`deps.py`), not scattered.

Q4. The repository pattern looks like indirection for its own sake — “my service could just call SQLAlchemy.” Defend it, and state when you would skip it.

Answer. The repository earns its keep four ways: (1) *testability* — services test against an in-memory fake, cutting suite time from minutes to seconds; (2) *query review surface* — all SQL lives in one layer, so performance and injection review is tractable; (3) *substitution* — caching, read-replicas, or a different store slot in behind the same interface; (4) *N+1 discipline* — eager-loading decisions are made once, in the layer that owns them. Skip it honestly when: the app is a small CRUD service whose “business logic” is the CRUD itself (the service + repo collapse into one thin layer), or you fully embrace SQLAlchemy as your repository abstraction and test against SQLite/testcontainers. The senior answer is cost-aware: patterns are bought with indirection; buy them when change and testing pressure justify the price — and be able to say where that line is.

Q5. Compare router-level dependencies, app-level dependencies, and middleware for enforcing authentication.

Answer. Middleware wraps every request at the ASGI level: it runs before routing, cannot use `Depends()`, sees raw request/response, and suits cross-cutting non-DI concerns (request IDs, timing, compression). App-level dependencies run inside the DI system for every route: they can chain other dependencies and appear in OpenAPI, but run after routing. Router-level dependencies scope enforcement to a subtree (`/admin/*`), which is usually the right granularity for authz. For authentication specifically: prefer dependencies, because (a) they integrate with OpenAPI security schemes so `/docs` shows the lock icon and supports “Authorize”; (b) handlers can receive the resolved user object; (c) per-route exceptions (public endpoints) are explicit. Use middleware only when you must act before routing or on responses. Senior trap to name: doing auth in middleware and *also* needing the user in handlers leads to stuffing state into `request.state` — workable, but you have rebuilt DI without its benefits.

Q6. What is a composition root, and why should `Depends()` wiring be centralized?

Answer. The composition root is the single location where abstract dependencies are bound to concrete implementations — in FastAPI, typically `api/deps.py` (plus test overrides). Centralizing it means: one place to read the object graph, one place to change a binding (swap `TaskRepository` for `CachedTaskRepository`), and one target for test overrides via `app.dependency_overrides`. Scattered wiring — routers constructing services inline, services importing concrete repos — recreates the coupling DI was meant to remove.

Q7. How does `dependency_overrides` relate to this chapter’s architecture?

Answer. `app.dependency_overrides[get_task_repo] = lambda: FakeRepo()` replaces a dependency for the whole app — the test seam DI creates. Because routers/service/repo are wired only through `Depends()`, tests can substitute any layer without monkeypatching module internals: fake repository for service-level tests, fake service for router contract tests, real everything against a test database for a few end-to-end paths. The architecture and the test strategy are the same decision viewed from two sides; Chapter 7 operationalizes this.

Q8. When does a class-based dependency beat a function-based one?

Answer. When the dependency has configuration or state worth grouping: a class with `__init__` parameters (e.g. `RoleChecker(["admin", "auditor"])`) produces parameterized, reusable dependency *instances* — the class’s `__call__` is the injectable. Also when the dependency bundles several related operations (a repository), or when you want `isinstance`-checkable types in

signatures for readability. Plain functions remain best for simple resource acquisition like `get_db`.

Q9. Five dependencies in one request all declare `Depends(get_settings)`. How many times does `get_settings` execute, and how do you reduce that to once per process?

Answer. Once per request: the per-request cache deduplicates within a single request's resolution. To make it once per *process*, cache at import/startup level — decorate with `@lru_cache` (the canonical pydantic-settings pattern: `@lru_cache def get_settings(): return Settings()`), or construct it in the lifespan handler and store on `app.state`. Settings are immutable configuration; re-reading the environment per request is pure waste. The same trick is wrong for sessions or anything stateful/mutable — those genuinely need per-request lifecycle.

Q10. Your codebase constructs services inline inside every handler. List the concrete problems this causes and outline the migration to `Depends`-based wiring.

Answer. Problems: sessions created outside the yield-dependency never close on exceptions (pool leaks); the wiring is duplicated per handler so a constructor change touches every endpoint (high change amplification); tests cannot substitute layers without monkeypatching; per-request caching is lost, so multiple sessions per request break transactional consistency. Migration: (1) introduce `get_db` as a yield dependency; (2) add provider functions in `deps.py` for repo and service; (3) mechanical sweep replacing inline construction with `Annotated[... , Depends(...)]` parameters; (4) add a lint forbidding `SessionLocal()` outside `deps.py`. Each step is independently shippable — no big-bang rewrite.

Database Integration with SQLAlchemy, SQLModel, and Alembic

APIs are stateless; their value lives in the database. This chapter builds a production persistence layer: SQLAlchemy models and sessions, SQLModel as an alternative, relationships and transactions, Alembic migrations, connection pooling, and the two performance topics that decide whether your API survives traffic — the N+1 query problem and pool sizing.

5.1 SQLite for Development, PostgreSQL for Production

SQLite is a zero-configuration, single-file engine — perfect for local development and tests. PostgreSQL is the default production choice: real concurrency (MVCC), rich types (JSONB, arrays, full-text, pgvector for Chapter 9), and operational maturity. The risk of developing on one and deploying on the other is behavioral drift: SQLite is loosely typed, ignores many constraints by default, and serializes writers. Mitigation: keep dev/test on SQLite only for fast unit layers, and run integration tests against real Postgres (testcontainers or CI services) — Chapter 7 shows how.

5.2 Engine, Sessions, and the Unit of Work

SQLAlchemy's two central objects: the **engine** (one per process; owns the connection pool) and the **session** (one per request; an identity-mapped *unit of work* that batches your changes and flushes them as SQL at commit). The session-per-request pattern is implemented as the yield dependency from Chapter 4 — now with transaction semantics:

```

1 from sqlalchemy import ForeignKey, String, create_engine, func
2 from sqlalchemy.orm import (DeclarativeBase, Mapped, mapped_column,
3                             relationship, sessionmaker, Session)
4
5 engine = create_engine(
6     settings.database_url,          # postgresql+psycopg://...
7     pool_size=10, max_overflow=5,   # pooling: see cost box
8     pool_pre_ping=True,             # detect dead connections
9     pool_recycle=1800,              # outlive no LB idle timeout
10 )
11 SessionLocal = sessionmaker(bind=engine, autoflush=False)
12

```

```

13 class Base(DeclarativeBase): ...
14
15 class User(Base):
16     __tablename__ = "users"
17     id: Mapped[int] = mapped_column(primary_key=True)
18     email: Mapped[str] = mapped_column(String(320), unique=True, index=True)
19     hashed_password: Mapped[str]
20     tasks: Mapped[list["Task"]] = relationship(back_populates="owner")
21
22 class Task(Base):
23     __tablename__ = "tasks"
24     id: Mapped[int] = mapped_column(primary_key=True)
25     title: Mapped[str] = mapped_column(String(200))
26     owner_id: Mapped[int] = mapped_column(
27         ForeignKey("users.id", ondelete="CASCADE"), index=True)
28     owner: Mapped[User] = relationship(back_populates="tasks")
29
30 # deps.py -- transaction per request: commit on success, rollback on error
31 def get_db():
32     db = SessionLocal()
33     try:
34         yield db
35         db.commit()
36     except Exception:
37         db.rollback()
38         raise
39     finally:
40         db.close()

```

Listing 5.1: Engine, models, and a transactional session dependency.

This dependency gives every endpoint atomic semantics: all writes in a request commit together or not at all. Services that need finer control (multi-step sagas, savepoints) manage `db.begin_nested()` explicitly.

5.2.1 CRUD Against Models

```

1 from sqlalchemy import select
2
3 class TaskRepository:
4     def __init__(self, db: Session):
5         self.db = db
6
7     def get(self, task_id: int) -> Task | None:
8         return self.db.get(Task, task_id)
9
10    def list_for_user(self, owner_id: int, limit: int, offset: int):
11        stmt = (select(Task)
12                .where(Task.owner_id == owner_id)
13                .order_by(Task.id)
14                .limit(limit).offset(offset))
15        return self.db.scalars(stmt).all()
16
17    def create(self, owner_id: int, data: TaskCreate) -> Task:
18        task = Task(owner_id=owner_id, **data.model_dump())

```

```

19         self.db.add(task)
20         self.db.flush()           # assigns task.id without committing
21         return task

```

Listing 5.2: Repository CRUD with SQLAlchemy 2.0-style queries.

5.3 SQLModel: One Class, Two Roles

SQLModel (by FastAPI’s author) merges Pydantic and SQLAlchemy: a class with `table=True` is simultaneously an ORM model and a Pydantic schema. For small services this removes the model/schema duplication; for larger ones, the Chapter 2 lesson still applies — storage and API contracts diverge, so most teams end up with separate `Create/Out` schemas anyway (SQLModel supports this with non-table classes). Treat SQLModel as ergonomic sugar over SQLAlchemy — it *is* SQLAlchemy underneath, so everything in this chapter (sessions, pooling, Alembic, N+1) applies unchanged.

5.4 Relationships, Lazy Loading, and the N+1 Problem

The default `relationship()` loading strategy is *lazy*: accessing `task.owner` triggers a separate `SELECT` at access time. Render a list of 100 tasks with their owners and you execute **1 + 100 queries** — the N+1 problem, the single most common ORM performance failure. The fix is explicit eager loading:

```

1  from sqlalchemy.orm import selectinload, joinedload
2
3  # 2 queries total: tasks, then owners via WHERE id IN (...)
4  stmt = (select(Task)
5          .options(selectinload(Task.owner))
6          .where(Task.owner_id.in_(team_ids))
7          .limit(100))
8
9  # 1 joined query -- better for to-one when you need a single roundtrip
10 stmt = select(Task).options(joinedload(Task.owner)).limit(100)

```

Listing 5.3: Killing N+1 with eager loading.

Rules of thumb: `selectinload` for collections (avoids row multiplication), `joinedload` for to-one references; in async code lazy loading raises errors anyway (no implicit I/O), which forces the good habit. Detection: log query counts per request in development and fail tests that exceed a budget (e.g. `assert query_count <= 5`).

5.5 Alembic Migrations

Production schemas change only through versioned, reviewed migrations. Alembic autogenerates diffs between your models and the live schema:

```

1  alembic init migrations
2  alembic revision --autogenerate -m "add tasks table"
3  alembic upgrade head           # apply (in CI/CD before app rollout)

```




Figure 5.1: The persistence topology: requests borrow pooled connections; schema changes flow through Alembic, never through the app.

```
4 alembic downgrade -1           # revert one step
```

Listing 5.4: The Alembic workflow.

Discipline that separates professionals: *always review autogenerated migrations* (autogenerate misses renames — it sees drop+add — and server defaults); write *backwards-compatible* migrations under zero-downtime deploys (add nullable column → deploy code writing both → backfill → add constraint → remove old column — the expand/contract pattern); and test *downgrade* paths before you need them at 3 a.m.

PRODUCTION CASE STUDY — The pool-exhaustion outage pattern (as seen at GitLab and many others)

A recurring postmortem shape across the industry, documented publicly by GitLab among others: traffic rises, p99 latency climbs, then the service hard-fails with `TimeoutError: QueuePool limit reached`. The chain: a slow query (often an unnoticed N+1 under new data volume) holds connections longer → pool saturates → requests queue waiting for connections → upstream timeouts fire → retries multiply load. GitLab’s remediations are the canonical checklist: add PgBouncer as a server-side pooler, set strict statement timeouts so no query can hold a connection indefinitely, instrument pool checkout wait time as a first-class metric, and fix the offending query patterns. The lesson: *the pool is a shared, finite resource whose exhaustion is a correlated failure* — one bad endpoint takes down every endpoint. Monitor checkout wait time, not just query latency.

COST MODEL & METRICS — Pool sizing arithmetic

By Little’s Law, the number of connections a service needs is

$$N_{conn} \approx \lambda \cdot \bar{t}_{db}$$

where λ = requests/s hitting the DB and $\bar{t}_{db}$ = mean DB-holding time per request. At 200 req/s × 25 ms = 5 connections — pools are smaller than intuition suggests. The global constraint:

$$W \cdot (\text{pool_size} + \text{max_overflow}) < \text{max_connections}_{PG} - C_{reserved}$$

with W = number of app processes. 10 pods × 4 workers × (10+5) = 600 demanded

connections against Postgres's default `max_connections=100`: an outage by configuration. Each Postgres connection also costs roughly 5–10 MB server RAM, so 600 connections $\approx$ 3–6 GB before any query runs — this is why PgBouncer (multiplexing thousands of client connections onto tens of server ones) appears in every serious Postgres deployment. Metrics: pool checkout wait p95, connections in use, and DB rows-examined per request.

SYSTEM DESIGN SCENARIO — Persistence for an order system: 500 writes/s, zero-downtime deploys

Requirements: 500 order-writes/s peak, strict atomicity (order + line items + inventory decrement), read-heavy reporting, schema evolves weekly, deploys must not drop traffic.

Reference approach: (1) PostgreSQL primary with the order-create path as a single transaction (service layer owns it; repository methods are transaction-agnostic). (2) Inventory decrement uses `SELECT ... FOR UPDATE` or an atomic `UPDATE ... WHERE qty >= n` to prevent oversell under concurrency. (3) Reporting reads go to a read replica; replica lag is acceptable for analytics, never for the checkout path — route by repository method, not globally. (4) PgBouncer in transaction mode in front of Postgres; app pools sized by Little's Law with headroom. (5) Weekly schema changes via expand/contract Alembic migrations gated in CI (migration runs against a production-like snapshot before rollout). (6) Indices reviewed with every new query; `EXPLAIN ANALYZE` output attached to PRs touching repositories.

5.6 Interview Questions

Q1. Explain the N+1 query problem, how to detect it in a FastAPI service, and the trade-offs between `selectinload` and `joinedload`.

Answer. N+1: fetching N parent rows, then lazily triggering one child query per row — 1+N round trips that scale with data volume, so it ships fine on a 10-row dev database and melts at 10K rows. Detection: per-request query counting (SQLAlchemy event hooks logging statement counts in dev), test assertions on query budgets, and in production a high ratio of DB-rows-examined to rows-returned or query-count-per-request metrics. Fixes: `selectinload` issues a second `WHERE id IN (...)` query — best for to-many (no row multiplication, two cacheable queries); `joinedload` does one `LEFT JOIN` — best for to-one (single round trip), but on to-many it multiplies rows (parent repeated per child) and inflates transfer. Senior addition: in async SQLAlchemy lazy loading raises at access time, converting silent performance bugs into loud errors — one argument for async persistence layers.

Q2. Walk through connection pool sizing for 8 pods × 4 workers against Postgres `max_connections=200`. What breaks first if you get it wrong, and where does PgBouncer fit?

Answer. Demand = $8 \times 4 \times (\text{pool_size} + \text{max_overflow})$. With the common default 5+10, that is 480 potential connections against 200 — under load, workers fail with connection errors, and unlike slow queries this failure is correlated across all endpoints. Size from Little's Law ($N \approx \lambda \bar{t}_{db}$ per worker plus headroom), e.g. `pool_size=4`, `max_overflow=2` $\Rightarrow$ 192 max, still tight — so introduce PgBouncer: app processes hold cheap connections to PgBouncer, which multiplexes them onto a small set of real server connections (transaction pooling mode). Costs of PgBouncer transaction mode worth naming: no session-level state (prepared statements, advisory locks, SET parameters) across transactions. Also set `pool_pre_ping` (dead-connection detection) and `pool_recycle` below any LB/NAT idle timeout. First-break symptom to monitor: pool checkout wait time, which rises before errors do.

Q3. Design the transaction boundary for “create order, decrement inventory, enqueue confirmation email.” Which parts belong in the DB transaction and which must not be in it?

Answer. Order insert and inventory decrement are one atomic unit: same transaction, with the decrement written to be concurrency-safe — either `UPDATE inventory SET qty = qty - :n WHERE sku = :sku AND qty >= :n` checking rowcount, or `SELECT ... FOR UPDATE` ordering locks consistently to avoid deadlock. The email *must not* be inside: a network call inside a transaction holds locks and a pooled connection for its full duration (latency amplification), and failure semantics are wrong — email failure should not roll back a paid order. Correct pattern: transactional outbox — insert an outbox row in the same transaction, commit, then a relay (Chapter 8's workers) delivers it with retries; or at minimum enqueue after commit (BackgroundTasks/Celery) accepting at-most-once. Senior signals: naming idempotency for the consumer side and exactly-once's impossibility — you choose at-least-once + idempotent handlers.

Q4. Your team deploys with zero downtime. Describe how to rename a column safely, and why `alembic revision -autogenerate` alone is insufficient.

Answer. Autogenerate diffs models against the schema and sees a rename as *drop + add* — destroying data; it also runs in one step, while zero-downtime requires old and new code to coexist against the same schema during rollout. Expand/contract sequence: (1) migration adds the new column, nullable; (2) deploy code that writes both columns, reads new-then-old; (3) backfill old values in batches (throttled, not in one locking UPDATE); (4) migration adds NOT NULL/constraints; (5) deploy code reading/writing only the new column; (6) migration drops the old column. Each step is individually revertible. Senior additions: review every autogenerated migration by hand; beware DDL that takes ACCESS EXCLUSIVE locks on large

tables (add column with volatile default on old PG, creating indexes without CONCURRENTLY); rehearse downgrade.

Q5. Compare SQLModel and plain SQLAlchemy + separate Pydantic schemas for a growing service.

Answer. SQLModel's pitch: one class is both table and schema, eliminating duplication and drift for simple CRUD; it is built on SQLAlchemy and Pydantic, so capabilities converge. The catch: storage and API contracts diverge in any nontrivial service (password hash stored but never returned; computed API fields not stored; versioned external schemas vs stable internal ones), so you reintroduce separate non-table models — recreating the very duplication, now within SQLModel. Plain SQLAlchemy 2.0 (Mapped typing) + explicit Pydantic Create/Update/Out schemas is more boilerplate but maximally explicit, with direct access to SQLAlchemy's full feature surface and documentation. Pragmatic recommendation: SQLModel for small services, prototypes, and teaching; explicit separation for multi-team or security-sensitive codebases — and note both share sessions/Alembic/pooling, so the choice is reversible.

Q6. What does the session-per-request pattern guarantee, and what goes wrong with a module-level global session?

Answer. Session-per-request (the yield dependency) gives each request an isolated unit of work: its own identity map, its own transaction, committed or rolled back exactly once, with the connection returned to the pool afterward. A global session shared across requests breaks all of it: concurrent requests interleave changes into one transaction (one request's error rolls back another's writes), the identity map grows unboundedly (memory leak), thread/loop-safety is violated, and a single poisoned transaction (failed statement, not rolled back) breaks every subsequent request with `InFailedSqlTransaction`.

Q7. Why is `pool_pre_ping` (or an equivalent) needed in cloud deployments?

Answer. Infrastructure between app and database — NAT gateways, load balancers, managed-DB proxies — silently drops idle TCP connections after timeouts the application never sees. A pooled connection can therefore be dead at checkout; the next query fails with a confusing operational error. `pool_pre_ping` issues a cheap liveness check at checkout and transparently replaces dead connections; `pool_recycle` proactively retires connections older than the infrastructure's idle limit. Together they convert sporadic “server closed the connection unexpectedly” incidents into a non-event.

Q8. What is the difference between `flush()` and `commit()`, and when do you need flush explicitly?

Answer. `flush()` emits pending SQL to the database within the open transaction — rows become visible to *this* transaction and server-generated values (autoincrement IDs, defaults) populate on the objects — but nothing is durable yet. `commit()` flushes, then commits the transaction, making changes durable and visible to others. Explicit flush is needed when you require a generated ID mid-transaction (e.g. create parent, use `parent.id` for children) while keeping the whole operation atomic under the request-level commit.

Q9. Reads are 20× writes and reporting queries are degrading checkout latency. Options?

Answer. Separate the workloads. (1) Read replica(s) for reporting/analytics traffic, routed at the repository layer (`ReportingRepo` binds to a replica engine); accept replica lag explicitly — never route read-your-own-writes paths (post-checkout order confirmation) to replicas. (2) Statement timeouts + a dedicated connection pool (or PgBouncer pool) for reporting so heavy queries cannot starve checkout connections. (3) Cache hot reads in Redis (Chapter 10) with explicit invalidation. (4) For genuinely analytical workloads, ship changes to a warehouse (CDC/ETL) and get them off the OLTP database entirely. Order matters: isolation (2) is an afternoon; replicas (1) are a sprint; warehouse (4) is a quarter.

Q10. How would you make repository-layer code testable without mocking SQLAlchemy internals?

Answer. Don't mock the ORM — test repositories against a real database: SQLite in-memory for speed where dialect differences don't matter, and a real Postgres (testcontainers, or CI service container) for anything using Postgres-specific behavior (JSONB, `FOR UPDATE`, constraint semantics). Wrap each test in a transaction rolled back at teardown for isolation and speed. Mock/fake at the *repository interface* when testing services (Chapter 4's layering), never inside it. Mocking `session.execute` return values tests your mocks, not your queries — the queries are precisely what repository tests exist to verify.

Authentication, Authorization, and API Security

Security is not a feature; it is a property of every endpoint you ship. This chapter builds the full authentication stack — password hashing, OAuth2 password flow, JWT access and refresh tokens, the current-user dependency, role-based access control, API keys — then hardens the perimeter: CORS, rate limiting, secure headers, secrets management, and the OWASP API risks that actually cause breaches.

6.1 Authentication vs Authorization

Authentication (authn) establishes *who you are*; **authorization** (authz) decides *what you may do*. The HTTP status codes encode the distinction: 401 means “I don’t know who you are” (missing/invalid credentials — retry with credentials); 403 means “I know exactly who you are, and no.” Most real-world breaches are authz failures, not authn failures — a logged-in user accessing *someone else’s* object. Keep the two phases architecturally separate: one dependency resolves identity; separate dependencies enforce permission.

6.2 Password Hashing

Passwords are never stored, encrypted, or logged — only *hashed* with an algorithm designed to be slow: bcrypt or argon2. Slowness is the feature: at ~100–300 ms per hash, offline cracking of a stolen database becomes economically painful, and unique per-password salts (built into both) defeat rainbow tables.

```

1 from datetime import datetime, timedelta, timezone
2 import jwt                                     # pyjwt
3 from fastapi import Depends, HTTPException, status
4 from fastapi.security import OAuth2PasswordBearer, OAuth2PasswordRequestForm
5 from passlib.context import CryptContext
6
7 pwd = CryptContext(schemes=["argon2"], deprecated="auto")
8 oauth2_scheme = OAuth2PasswordBearer(tokenUrl="/api/v1/auth/token")
9
10 def hash_password(p: str) -> str:             return pwd.hash(p)
11 def verify_password(p: str, h: str) -> bool:   return pwd.verify(p, h)

```

```

12
13 def create_token(sub: str, ttl: timedelta, kind: str) -> str:
14     now = datetime.now(timezone.utc)
15     payload = {"sub": sub, "type": kind, "iat": now, "exp": now + ttl,
16               "iss": "api.example.com", "aud": "example-clients"}
17     return jwt.encode(payload, settings.jwt_secret, algorithm="HS256")
18
19 @router.post("/token")
20 def login(form: OAuth2PasswordRequestForm = Depends(), db: DB = None):
21     user = db.scalar(select(User).where(User.email == form.username))
22     if user is None or not verify_password(form.password,
23                                           user.hashed_password):
24         # same error either way: no account enumeration
25         raise HTTPException(401, "Incorrect email or password")
26     return {
27         "access_token": create_token(str(user.id),
28                                     timedelta(minutes=15), "access"),
29         "refresh_token": create_token(str(user.id),
30                                     timedelta(days=14), "refresh"),
31         "token_type": "bearer",
32     }

```

Listing 6.1: Password hashing and the OAuth2 login endpoint issuing JWTs.

6.3 JWT Access Tokens and the Current-User Dependency

A JWT is a signed (not encrypted!) triplet `header.payload.signature`: the server can verify integrity without a database lookup, which is what makes JWTs scale — any worker validates any token statelessly. The payload is readable by anyone; never put secrets in it. Verification *must* pin the algorithm and validate `exp`, `iss`, and `aud`:

```

1 def get_current_user(
2     token: Annotated[str, Depends(oauth2_scheme)], db: DB,
3 ) -> User:
4     try:
5         payload = jwt.decode(
6             token, settings.jwt_secret,
7             algorithms=["HS256"], # pin: never trust header
8             issuer="api.example.com", audience="example-clients",
9         )
10    except jwt.InvalidTokenError:
11        raise HTTPException(401, "Invalid or expired token",
12                            headers={"WWW-Authenticate": "Bearer"})
13    if payload.get("type") != "access":
14        raise HTTPException(401, "Wrong token type")
15    user = db.get(User, int(payload["sub"]))
16    if user is None or not user.is_active:
17        raise HTTPException(401, "User inactive or deleted")
18    return user
19
20 CurrentUser = Annotated[User, Depends(get_current_user)]
21
22 class RequireRole: # parameterized authz dependency
23     def __init__(self, *roles: str): self.roles = set(roles)

```

```

24     def __call__(self, user: CurrentUser) -> User:
25         if user.role not in self.roles:
26             raise HTTPException(403, "Insufficient permissions")
27         return user
28
29 admin_router = APIRouter(
30     prefix="/admin",
31     dependencies=[Depends(RequireRole("admin"))],    # router-wide RBAC
32 )

```

Listing 6.2: The current-user dependency chain with role-based access control.

Refresh tokens solve the stateless-revocation dilemma: access tokens stay short-lived (5–15 min) and stateless; refresh tokens are long-lived, stored server-side (hashed, like passwords), and exchanged for new access tokens — so they *can* be revoked, and rotation with reuse-detection (a presented-twice refresh token revokes the whole family) contains theft. Object-level authorization — the check that `task.owner_id == user.id` — belongs in the service layer, because it needs the loaded object; role checks belong in dependencies.




Figure 6.1: Defense in depth: every request crosses transport, throttling, browser-policy, authentication, and authorization layers before touching a handler.

6.4 API Keys, CORS, Rate Limiting, and Headers

API keys suit machine-to-machine callers: generate high-entropy tokens, store only their hashes, accept them via `X-API-Key` (a `Security(APIKeyHeader(...))` dependency), scope them to specific permissions, and support rotation by allowing two active keys per client.

CORS is a browser-enforced policy controlling which web origins may read your responses. Configure `CORSMiddleware` with an explicit origin list; combining wildcard origins with credentials (`allow_origins=["*"]` plus credentials enabled) is forbidden by the spec and a misconfiguration classic. CORS protects browser users; it does nothing against curl — it is not an access-control mechanism.

Rate limiting bounds abuse and cost. The token-bucket algorithm is the workhorse: each identity gets a bucket of B tokens refilling at r/s ; a request consumes one; empty bucket $\Rightarrow$ 429 with `Retry-After`. Enforce at the gateway for coarse limits and in-app (e.g. `slowapi`, Redis-backed so limits are global across workers) for per-user/per-endpoint nuance — login and signup endpoints get the strictest budgets.

Secure headers: `Strict-Transport-Security`, `X-Content-Type-Options: nosniff`, `X-Frame-Options: DENY`, a restrictive `Content-Security-Policy` for any HTML you serve, and `Cache-Control: no-store` on authenticated responses. A tiny middleware sets them app-wide.

Secrets never live in code or images: inject via environment from a secrets manager (Vault, AWS Secrets Manager), load through `pydantic-settings` (`SecretStr` fields), rotate regularly, and scan repos for leaks in CI. A leaked `JWT_SECRET` means an attacker mints arbitrary identities — treat it like a root password.

6.5 OWASP API Security: Where Breaches Actually Happen

The OWASP API Security Top 10's perennial #1 is **BOLA** — Broken Object-Level Authorization: `GET /orders/{id}` checks that you are logged in but not that order `id` is *yours*; an attacker iterates IDs and reads everyone's data. Defenses: ownership checks in the service layer on every object access (filter queries by `owner_id` rather than checking after load), non-enumerable UUIDs as a speed bump (not a fix), and authz tests as first-class test cases (Chapter 7). The rest of the list maps to this chapter's sections: broken authentication (weak hashing, unpinned JWT algorithms), unrestricted resource consumption (no rate limits, unbounded `limit` params — Chapter 3), broken function-level authorization (missing role checks on admin routes), and security misconfiguration (permissive CORS, verbose errors).

PRODUCTION CASE STUDY — Optus 2022 — an unauthenticated API and 9.8M records

The 2022 Optus breach, one of Australia's largest, reportedly involved an exposed, internet-reachable API endpoint that required *no authentication* and used *sequential, enumerable customer identifiers* — the attacker simply iterated IDs and harvested ~9.8 million customers' personal data (passports, licenses, addresses). Estimated costs ran well past \$100M including a \$140M (AUD) customer-remediation program, regulatory action, and churn. Every layer of this chapter would have stopped or blunted it: any authentication at all; object-level authorization; non-enumerable identifiers; rate limiting that would flag millions of sequential requests; and egress monitoring. The meta-lesson: forgotten endpoints (test/legacy hosts) carry the same blast radius as your main API — inventory and uniformly enforce policy on *everything* routable.

COST MODEL & METRICS — The economics of credential attacks

Credential-stuffing cost to an attacker is

$$C_{\text{attack}} = \frac{N_{\text{attempts}}}{R_{\text{allowed}}} \cdot C_{\text{infra}}$$

— driven by your rate limit R_{allowed} , which is the only variable you control. Defense-side, expected breach cost is

$$E[C_{\text{breach}}] = p_{\text{success}} \cdot (C_{\text{records}} + C_{\text{regulatory}} + C_{\text{churn}}),$$

with industry figures (IBM Cost of a Data Breach) averaging \$165 per leaked record and \$4.5M per breach. Hashing economics: at argon2's ~200 ms/hash, cracking one strong password from a stolen table costs the attacker $\sim 10^{10} \times$ more compute than MD5 would. Metrics to alert on: failed-login rate (per IP and global), 401/403 ratios per endpoint, token-reuse detections, and impossible-travel logins.

SYSTEM DESIGN SCENARIO — Auth for a B2B SaaS: browsers, mobile, and partner servers

Requirements: web SPA + mobile app for humans, server-to-server partner integrations, tenant isolation, SOC 2 audit logging, 50K users.

Reference approach: (1) Humans: OAuth2 password (or authorization-code via an IdP) flow issuing 15-min JWT access tokens + rotating refresh tokens; web stores refresh tokens in httpOnly Secure SameSite cookies (XSS-resistant), mobile in the platform keystore. (2) Partners: hashed API keys with per-key scopes and rate limits; optionally mTLS for high-value partners. (3) Tenant isolation: every JWT carries `tenant_id`; repositories filter by it unconditionally (a base-repository pattern makes it impossible to forget); cross-tenant access is a 404, not a 403, to avoid existence leaks. (4) RBAC: roles in the token for coarse gates, ownership checks in services for object-level. (5) Audit: an append-only auth-events table (login, refresh, revoke, permission change) with request IDs — satisfying SOC 2 and powering anomaly detection. (6) Rate limits per user, per key, and per IP at the gateway plus Redis-backed app limits on login/refresh.

6.6 Interview Questions

Q1. Why must JWT verification pin the algorithm, and what is the `alg=none` attack?

Answer. The JWT header declares which algorithm signed it — and the header is attacker-controlled. Historic library defaults honored it: an attacker setting `"alg": "none"` (signature-less JWTs, valid per spec) could submit unsigned tokens that verified successfully. A second variant: against RS256 services, declare HS256 and sign with the *public* key as the HMAC secret — libraries that pass the public key as the verification key validated it. Defenses: always pass an explicit `algorithms=["HS256"]` allowlist to `decode()`; never derive the algorithm from the token; validate `exp`, `iss`, `aud`; and use modern libraries that reject `none` by default. Senior framing: this is an instance of the general rule — *never let untrusted input select the security mechanism that validates it*.

Q2. Design a revocation story for stateless JWTs. Walk through the trade-offs.

Answer. Stateless verification means no per-request DB check — which is exactly why you cannot revoke an access token once issued. Options: (1) *Short TTL + refresh rotation* (the standard): access tokens live 5–15 min, so revocation latency is bounded by TTL; refresh tokens are stored server-side (hashed), revocable, and rotated on every use with reuse-detection revoking the token family on replay. (2) *Denylist*: revoked token IDs (`jti`) in Redis until expiry — restores instant revocation at the cost of a cache lookup per request (still far cheaper than DB). (3) *Versioned claims*: a `token_version` on the user row, bumped to invalidate all outstanding tokens, checked on sensitive routes only. Senior answer: combine them — short TTL globally, denylist for logout/compromise events, version bump for password change — and state the latency/consistency budget that justifies each.

Q3. What is BOLA, why does it top the OWASP API list, and how do you make it structurally impossible rather than reviewer-dependent?

Answer. Broken Object-Level Authorization: authenticated users accessing objects they don't own because the endpoint checks *identity* but not *ownership* — GET /orders/123 with someone else's ID. It tops the list because it is easy to introduce (every new endpoint is a new opportunity), easy to exploit (iterate IDs), and invisible to schema-level tooling. Structural defenses: (1) query-scoping — repository methods take the caller's identity and filter (WHERE owner_id = :uid) so unauthorized rows are never loaded; expose no unscoped get(id) on tenant-owned resources; (2) a base repository/dependency that injects tenant/owner filters automatically; (3) authz test suites that fetch every owned resource as a different user and assert 404; (4) non-enumerable IDs (UUIDv4) as friction, never as the control. Return 404 not 403 for unowned objects to avoid confirming existence.

Q4. Compare storing web-client tokens in localStorage vs httpOnly cookies, including the CSRF implications.

Answer. localStorage: readable by JavaScript, so any XSS exfiltrates tokens — and modern SPAs have large dependency-supply-chain XSS surfaces; no CSRF risk (the browser never attaches it automatically). httpOnly Secure SameSite cookie: invisible to JS (XSS cannot read it — though XSS can still *use* the session in-page), but auto-attached by the browser, reintroducing CSRF — mitigated by SameSite=Lax/Strict, CSRF tokens for state-changing requests, and origin checks. Standard senior position: refresh token in an httpOnly cookie scoped to the refresh path; access token in memory only (never persisted), short-lived; this bounds both XSS (no persistent token to steal) and CSRF (refresh endpoint validates origin; API requires the Authorization header CSRF cannot forge). Mobile: platform keystores, no cookies.

Q5. Why bcrypt/argon2 instead of SHA-256 for passwords, and how do you choose/maintain cost parameters in production?

Answer. SHA-256 is designed to be *fast* — billions of hashes/s on GPUs — so a stolen table falls to offline brute force regardless of salting. bcrypt/argon2 are deliberately expensive and tunable; argon2 is additionally memory-hard, neutering GPU/ASIC advantages. Parameter policy: target 100–300 ms per hash on *your production hardware* — slow enough to hurt attackers, fast enough that login throughput and your event loop survive (hash in a threadpool; it is CPU-bound!). Revisit annually as hardware improves; passlib's deprecated="auto" transparently re-hashes legacy-cost passwords at next successful login, so upgrades need no mass reset. Senior additions: rate-limit login attempts (online attacks bypass hash cost), and pepper (server-side secret added to hashing) as defense against DB-only compromises.

Q6. A login endpoint returns “email not found” vs “wrong password.” What is wrong, and what else leaks account existence?

Answer. Differentiated errors enable *account enumeration*: attackers validate harvested email lists for targeted phishing, credential stuffing, and social engineering. Return one generic message and identical status code for both failures. Other enumeration channels to close: signup (“email already registered” — respond “check your email” uniformly), password reset (always “if the account exists, we sent a link”), and timing — verifying a password takes 200 ms but the missing-user branch returns instantly, so hash a dummy password in that branch to equalize response time.

Q7. What does CORS actually protect, and why is it not an access-control mechanism?

Answer. CORS is enforced by *browsers*, governing whether `evil.com`’s in-page JavaScript may read responses from your API using a victim’s ambient credentials. It protects users from cross-origin data theft in browser contexts — nothing else. `curl`, mobile apps, and server-side attackers ignore it entirely: the server happily processes the request regardless of CORS configuration. Therefore CORS misconfiguration matters (it can expose authenticated browser sessions), but authentication/authorization must never assume CORS blocked anything. And `allow_origins=["*"]` with credentials is both spec-forbidden and the classic misconfig: with reflected origins it hands any site full API access as the victim.

Q8. Where should rate limiting live — gateway, middleware, or dependency — and what does Redis add?

Answer. Layered: coarse anonymous limits (per-IP) at the gateway/CDN edge where requests are cheapest to reject; identity-aware limits (per-user, per-API-key, per-endpoint — strictest on login/signup/reset) in-app via middleware or dependency where you know who the caller is. Redis makes in-app limits *global*: with N workers, in-memory counters give each worker its own budget ($N \times$ the intended limit, and inconsistent 429s); a Redis token bucket (atomic via Lua/INCR-EXPIRE) is shared, consistent state across all workers and pods. Always return `Retry-After` so well-behaved clients back off rather than hammer.

Q9. Your JWT secret leaked in a public repo. Walk through the incident response.

Answer. (1) *Rotate immediately*: deploy a new secret; to avoid mass logout, briefly accept-old/issue-new with a dual-key verification window (or accept the forced re-login — often the right call for a confirmed leak). (2) *Invalidate*: bump server-side token versions / clear refresh-token store so attacker-minted tokens die with the old key. (3) *Purge*: rewrite git history

or rotate-and-retire the repo copy; revoke any other secrets in the same blast radius. (4) *Audit*: search logs for anomalous token usage since the leak window — attacker-minted JWTs verify as legitimate, so hunt by behavior (impossible travel, unusual scopes, enumeration patterns). (5) *Prevent*: secrets manager + env injection, pre-commit and CI secret scanning, and key rotation as a rehearsed runbook, not an improvisation. Asymmetric signing (RS256) with the private key in a KMS/HSM reduces this entire class: services hold only the public key.

Q10. When do API keys beat OAuth2/JWT for authentication, and what hygiene do they require?

Answer. API keys fit machine-to-machine callers with long-lived, non-interactive access: partner backends, CI systems, webhooks — no user consent flow, no token refresh choreography, trivially debuggable. They lose to OAuth2/JWT when you need delegated user identity, scopes that vary per session, or third-party authorization. Hygiene: generate with CSPRNG entropy (≥ 128 bits), show once, store only hashes; prefix keys (`sk_live_...`) so scanners can find leaks; scope each key to least privilege; per-key rate limits and usage metering; support two concurrent active keys for zero-downtime rotation; and log key-id (never the key) on every authenticated request for audit and revocation targeting.

Testing FastAPI Applications

Tests are the only scalable way to know your API works — and in an API, “works” includes contracts, status codes, validation, and authorization, not just happy paths. This chapter covers pytest and TestClient mechanics, dependency overrides (the payoff of Chapter 4), test databases, auth testing, CI workflows, and how to think about coverage honestly.

7.1 Why API Testing Is Different

An API’s regression surface is its *contract*: URLs, schemas, status codes, error envelopes, and authorization rules that external consumers depend on. A refactor that keeps your unit tests green but changes a response field name is a production incident for every client. So API test suites assert at two levels: **unit tests** verify business logic in isolation (fast, no I/O), and **integration tests** verify the assembled stack — routing, validation, serialization, database — behaves as the contract promises. The classic pyramid applies: many unit tests, fewer integration tests, a handful of end-to-end smoke tests.

7.2 pytest and TestClient

FastAPI’s `TestClient` (built on `httpx`) runs your ASGI app *in-process* — no server, no network, no ports — while exercising the full middleware/routing/validation stack:

```
1 # tests/conftest.py
2 import pytest
3 from fastapi.testclient import TestClient
4 from sqlalchemy import create_engine
5 from sqlalchemy.orm import sessionmaker
6 from sqlalchemy.pool import StaticPool
7
8 from app.main import app
9 from app.api.deps import get_db
10 from app.models import Base
11
12 engine = create_engine(
13     "sqlite://", # in-memory
14     connect_args={"check_same_thread": False},
15     poolclass=StaticPool, # one shared connection
16 )
```

```

17 TestingSession = sessionmaker(bind=engine)
18
19 @pytest.fixture()
20 def db():
21     Base.metadata.create_all(engine)
22     session = TestingSession()
23     try:
24         yield session
25     finally:
26         session.rollback()
27         session.close()
28         Base.metadata.drop_all(engine)    # clean slate per test
29
30 @pytest.fixture()
31 def client(db):
32     app.dependency_overrides[get_db] = lambda: db
33     with TestClient(app) as c:          # runs lifespan events
34         yield c
35     app.dependency_overrides.clear()

```

Listing 7.1: Foundation: fixtures, a test database, and dependency overrides.

The single most important line is `app.dependency_overrides[get_db]`: because every endpoint obtains its session through DI, one override redirects the *entire application* to the test database — no monkeypatching, no environment tricks. This is the architectural payoff promised in Chapter 4.

7.3 Testing CRUD, Validation, and Contracts

```

1 def test_create_task_returns_201_and_public_schema(client):
2     resp = client.post("/api/v1/tasks",
3                        json={"title": "write tests"})
4     assert resp.status_code == 201
5     body = resp.json()
6     assert set(body) == {"id", "title", "description", "status"}
7     # leak check: internal fields must never appear
8     assert "owner_id" not in body
9
10 def test_get_missing_task_is_404(client):
11     assert client.get("/api/v1/tasks/9999").status_code == 404
12
13 def test_title_too_long_is_422_with_field_location(client):
14     resp = client.post("/api/v1/tasks", json={"title": "x" * 201})
15     assert resp.status_code == 422
16     err = resp.json()["detail"][0]
17     assert err["loc"] == ["body", "title"]    # machine-readable location
18
19 @pytest.mark.parametrize("payload", [
20     {},                                # missing required field
21     {"title": ""},                     # violates min_length
22     {"title": "ok", "status": "bogus"}, # invalid enum
23 ])
24 def test_invalid_payloads_rejected(client, payload):
25     assert client.post("/api/v1/tasks", json=payload).status_code == 422

```

Listing 7.2: Contract tests: status codes, schemas, validation errors, and leak checks.

Note what these tests pin down: the *shape* of success responses (including that sensitive fields are absent — a security test disguised as a schema test), the status code for absence, and the structure of validation errors that clients parse programmatically.

7.4 Testing Authentication and Authorization

Auth tests are where API testing earns its keep, because authz bugs (BOLA, Chapter 6) are invisible to schema tooling:

```

1 from app.api.deps import get_current_user
2
3 def as_user(user):
4     """Helper: impersonate a user by overriding the auth dependency."""
5     app.dependency_overrides[get_current_user] = lambda: user
6
7 def test_anonymous_request_is_401(client):
8     app.dependency_overrides.pop(get_current_user, None)
9     assert client.get("/api/v1/me").status_code == 401
10
11 def test_user_cannot_read_anothers_task(client, db):
12     alice, bob = make_user(db, "alice"), make_user(db, "bob")
13     task = make_task(db, owner=alice)
14     as_user(bob)
15     # 404, not 403: do not confirm existence (Chapter 6)
16     assert client.get(f"/api/v1/tasks/{task.id}").status_code == 404
17
18 def test_admin_route_rejects_non_admin(client, db):
19     as_user(make_user(db, "carol", role="viewer"))
20     assert client.delete("/api/v1/users/1").status_code == 403
21
22 def test_expired_token_is_401_with_www_authenticate(client):
23     token = create_token("1", timedelta(minutes=-1), "access")
24     resp = client.get("/api/v1/me",
25                       headers={"Authorization": f"Bearer {token}"})
26     assert resp.status_code == 401
27     assert resp.headers["WWW-Authenticate"] == "Bearer "
```

Listing 7.3: Auth tests: identity via override, authorization as a first-class case.

Two styles appear here, both legitimate: *overriding* `get_current_user` (fast; tests authorization logic, skipping token mechanics) and *real tokens* (slower; tests the token pipeline itself). Use overrides for the many authz cases, real tokens for the few authn cases.

7.4.1 Mocking External Dependencies

Anything that leaves your process — payment gateways, LLM APIs, email — gets faked at the boundary you own: override the service-level dependency (`get_llm_client`) with a stub returning canned responses, or use `respx/httpx` mock transports to intercept HTTP. Never let test suites call

real third-party APIs: they are slow, flaky, rate-limited, and cost money — and they make CI failures uninvestigable.




Figure 7.1: The test topology: TestClient drives the real stack in-process; dependency overrides swap externalities for fast, deterministic fakes.

7.5 Test Databases and CI Workflow

The pragmatic two-tier strategy: **SQLite in-memory** for the bulk of tests (milliseconds, zero setup) and **real PostgreSQL** for anything touching Postgres-specific behavior — JSONB, FOR UPDATE, constraint semantics, migrations. In CI, Postgres runs as a service container; locally, testcontainers spins it up on demand. Run Alembic migrations against the CI Postgres rather than `create_all()` — it tests the migration chain itself, catching the broken-merge-migration problem before deploy.

A typical GitHub Actions pipeline:

```

1 jobs:
2   test:
3     runs-on: ubuntu-latest
4     services:
5       postgres:
6         image: postgres:16
7         env: {POSTGRES_PASSWORD: test}
8         ports: ["5432:5432"]
9     steps:
10      - uses: actions/checkout@v4
11      - run: pip install -e ".[dev]"
12      - run: ruff check . && mypy app
13      - run: alembic upgrade head # migrations are tested too
14      - run: pytest --cov=app --cov-fail-under=85 -q
  
```

Listing 7.4: CI workflow: lint, types, tests with coverage gate.

7.6 Code Coverage, Honestly

Coverage measures *execution*, not *verification*: a test that calls an endpoint and asserts nothing scores the same coverage as one that pins the whole contract. Use coverage as a *floor and a flashlight* — a gate (e.g. 85%) to catch untested-by-accident code, and the per-file report to find dark corners (error handlers, except branches) — never as a quality target in itself. Branch coverage (`-cov-branch`) is stricter and more meaningful than line coverage for API code, where the interesting bugs live in `if`

arms.

PRODUCTION CASE STUDY — The 25-minute test suite that stopped being run

A pattern observed across many teams (and documented in countless postmortems): an e-commerce startup's suite ran every test against a shared, Docker-composed Postgres with full table truncation between tests — 25 minutes per run. Engineers responded rationally: they ran only “their” tests locally and pushed, letting CI find the rest; then they began merging on red “because it’s probably that flaky cart test.” Within a quarter, a checkout-tax regression shipped to production — covered by a test that was red along with 14 known-flaky others. The rebuild followed this chapter: 80% of tests moved to SQLite-in-memory with dependency overrides (suite: 90 seconds), Postgres-specific tests isolated behind a marker (`-m pg`, 3 minutes, parallel CI job), flaky tests quarantined with an owner and a deadline, and merge-on-red disabled. The metric that mattered was not coverage — it was *suite latency*, because suite latency determines whether tests are actually run.

COST MODEL & METRICS — The economics of test latency

Let n = test runs per engineer-day, T = suite wall-clock time, E = engineers, c = loaded cost per engineer-hour. Direct waiting cost:

$$C_{wait} = E \cdot n \cdot T \cdot c \quad (8 \text{ eng} \times 10 \times 25 \text{ min} \times \$100/h \approx \$3,300/\text{day}).$$

The indirect cost dominates: defect escape probability rises as runs become rarer ($n \propto 1/T$ empirically), and a shipped defect costs $C_{prod} \approx 10\text{--}100\times$ a caught-in-CI one. Optimizing T from 25 min to 2 min pays for itself in days. Track: suite p50 latency, flake rate (reruns-to-green), time-to-green per PR, and escaped-defect count per quarter.

SYSTEM DESIGN SCENARIO — Test strategy for a payments API (compliance-grade)

Requirements: card-charging endpoints, idempotency keys, PCI-adjacent audit requirements, 15 engineers, deploys $10\times/\text{day}$; regulators ask for evidence of authorization testing.

Reference approach: (1) Pyramid: service-layer unit tests with fake repositories (majority); TestClient integration tests for every endpoint's contract (status codes, schemas, error envelopes); a thin staging smoke suite against the real payment sandbox. (2) Authz matrix as data: a parametrized test iterating (role $\times$ endpoint $\times$ ownership) asserting allowed/denied — the artifact you hand auditors. (3) Idempotency tests: same key replayed concurrently (threads) must yield one charge — this catches the race, not just the happy path. (4) Payment provider faked via httpx mock transport with recorded fixtures; a nightly contract-verification job replays fixtures against the sandbox to detect provider drift. (5) Migrations tested in CI against a production-schema snapshot. (6) Coverage gate 90% on the payments module specifically, with branch coverage; mutation testing (`mutmut`) quarterly to audit assertion strength.

7.7 Interview Questions

Q1. How do FastAPI dependency overrides work, and why are they superior to monkeypatching for API tests?

Answer. `app.dependency_overrides` is a dict mapping original dependency callables to replacements; during resolution FastAPI checks it before invoking the real dependency, so `overrides[get_db] = lambda: test_session` redirects every endpoint's session sourcing in one line. Superiority over monkeypatching: (1) it operates at the *declared interface* — the same seam production wiring uses — so tests don't depend on module paths or import order (`patch("app.api.tasks.SessionLocal")` breaks when someone moves an import; overrides don't); (2) it is scoped and discoverable — one dict to inspect, cleared in fixture teardown; (3) it composes — override auth and DB and an LLM client independently. Senior caveats: overrides bypass the real dependency entirely (its validation logic goes untested — cover it separately), and forgetting `overrides.clear()` leaks state between tests; manage it in fixtures.

Q2. Design the test-database strategy for a team using Postgres in production. Defend SQLite-in-memory against “test what you fly.”

Answer. Two tiers. Tier 1 (default): SQLite in-memory with `StaticPool` — milliseconds per test, perfect isolation via create/drop per test, and it covers the 80–90% of tests exercising business logic, contracts, and serialization where the database is incidental. Tier 2 (marked): real Postgres via testcontainers/CI service for dialect-sensitive behavior — JSONB operators, row locking (`FOR UPDATE`), deferred constraints, sequences, migration replay. “Test what you fly” is answered by *routing* tests to the right tier, not by paying Postgres latency on every assertion: the failure modes SQLite masks are enumerable (typing laxity, concurrency, dialect SQL), so anything touching them is tier 2 by policy. Senior additions: run Alembic migrations (not `create_all`) in tier 2 so the migration chain itself is tested; wrap tier-2 tests in rolled-back transactions for speed; and keep tier ratios visible — if tier 2 grows past ~20%, the repository layer is leaking dialect concerns upward.

Q3. What belongs in a unit test vs an integration test for a FastAPI endpoint, and how does the Chapter-4 architecture determine this?

Answer. Unit tests target the service layer directly: business rules, edge cases, domain exceptions — constructed with fake repositories, no HTTP, no DB; they are numerous and sub-millisecond. Integration tests go through `TestClient` and assert the *contract*: routing, validation (422 shapes), serialization (`response_model` filtering), status-code mapping of domain exceptions, auth enforcement — with a real test DB when persistence semantics matter. The layered architecture is what makes the split possible: thin routers mean integration tests need only cover translation concerns (few per endpoint), while logic lives where unit tests reach it cheaply. Anti-pattern to

name: logic in handlers forces every business edge case through HTTP round-trips — the suite gets slow, so it stops being run. The test pyramid is downstream of architecture.

Q4. How would you test that an endpoint never leaks sensitive fields, in a way that survives future refactors?

Answer. Layered: (1) *Exact-shape assertions* — integration tests asserting `set(resp.json()) == {expected fields}`, which fail on *additions*, unlike "password" not in body checks that only catch known names. (2) *Schema-driven generation* — derive the expected key set from the public Pydantic model so the test tracks intentional contract changes automatically. (3) *Static gate* — a lint/CI check that every route declares `response_model` (the structural defense from Chapter 2). (4) *Canary middleware* in staging scanning responses for sensitive key patterns (*password*, *secret*, token formats) as a last net. The senior point: (1)+(2) test the contract, (3) makes the vulnerable configuration unrepresentable, (4) catches what design review missed — defense in depth applied to testing.

Q5. Your CI suite has a 3% flake rate across 2,000 tests. Quantify the damage and lay out a remediation plan.

Answer. Damage: $P(\text{green} \mid \text{all-correct}) = 0.97^k$ for k flaky tests touched — at suite level a 3% flake rate means most runs contain a false failure, so engineers retry (CI cost doubles, latency doubles) and — worse — learn to distrust red, which is how real regressions ship (alarm fatigue). Remediation: (1) *measure* — rerun-to-green telemetry identifying flaky tests by name; (2) *quarantine* — move them to a non-blocking job with an owner and a fix-by date (never delete silently — they may guard real behavior); (3) *fix by class*: shared state between tests (fixture leakage, unclear `dependency_overrides`), time dependence (freeze time), ordering assumptions (run with `-p random_order`), real network calls (mock them), and async race conditions (await completion, don't sleep); (4) *prevent* — new tests run 10× in a PR gate before joining the suite. Track flake rate as an SLO; 3% is an incident, not a nuisance.

Q6. Why does TestClient not require a running server, and what are the limits of in-process testing?

Answer. TestClient speaks ASGI directly: it constructs the scope/receive/send triple and calls your app as a coroutine in-process — exercising middleware, routing, validation, and handlers without sockets. Limits: it does not test the ASGI *server* (Uvicorn config, worker counts, timeouts), TLS, reverse-proxy behavior (header forwarding, body size caps), HTTP/2, real network latency, or cross-process concerns (pool sizing across workers). Those belong to a thin staging smoke layer. Knowing the boundary — contract in-process, infrastructure in staging — is the practical skill.

Q7. How do you test code that depends on the current time (token expiry, rate limits)?

Answer. Make time injectable: a `get_now()` dependency or a clock parameter on services, overridden in tests — or use `freezegun/time-machine` to freeze module-level clocks. Then expiry tests become deterministic: issue a token at t_0 , advance the clock past TTL, assert 401 — no `sleep()`. Sleeping in tests is the cardinal sin: it is slow *and* flaky (CI machines stall). Senior detail: JWT libraries validate `exp` against real time, so either freeze time globally for the verification call or mint tokens with already-past expiry (negative TTL) as this chapter's example does.

Q8. What does `with TestClient(app) as client` do that bare `TestClient(app)` does not?

Answer. The context-manager form runs the application's lifespan events: startup before the first request (model loading, pool creation, `app.state` initialization) and shutdown on exit. Bare construction skips them — fine for pure routing tests, but anything reading `app.state` or resources initialized at startup will fail or silently use stale state. Rule: if production behavior depends on lifespan (Chapter 9's model loading does), test through the context manager.

Q9. Why is asserting on the 422 error body (not just the status code) worth the test brittleness?

Answer. Because the error body is part of the public contract: clients parse `detail[] .loc` and `detail[] .type` to map errors to form fields and to localize messages. A Pydantic major-version upgrade or a custom exception handler can change this structure while every status-code-only test stays green — breaking every consumer. Pin the parts clients rely on (`loc` paths, `type` codes), not the human-readable `msg` strings, and you get contract protection without brittleness to wording changes.

Q10. Your coverage is 95% but production incidents keep coming from “tested” code. Diagnose.

Answer. Coverage measures execution, not assertion strength. Likely causes: tests that call code but assert weakly (status code only, no body); happy-path bias — the 5% uncovered is precisely the error handling where incidents live (check *branch* coverage, not line); incidents arising in untestable-in-CI dimensions — concurrency, pool exhaustion, real-data shapes, third-party drift; and assertion rot after refactors (tests updated to pass rather than to verify). Remedies: mutation testing (`mutmut/cosmic-ray`) to score whether tests *fail* when code is wrongly changed; require a regression test per incident (the postmortem rule); branch-coverage gates on error-handling modules; and contract tests against recorded real-world payloads, not just hand-built minimal ones.

Async Programming, Background Tasks, and External Services

Concurrency is FastAPI's superpower and its sharpest edge. This chapter builds the complete mental model — `def` vs `async def`, blocking vs non-blocking, the event loop — then applies it to real integration work: `httpx` clients with timeouts and retries, background tasks, Celery workers, webhooks, file uploads, streaming responses, and WebSockets.

8.1 The Event Loop Mental Model

An `asyncio` program runs one loop on one thread. Coroutines (`async def`) run until they `await` something not ready — then they *yield control*, and the loop runs another coroutine. The result: thousands of concurrent I/O waits on one thread, with no locks and no thread-switch overhead. The contract has one clause: **never block the loop**. A single synchronous `time.sleep(5)`, `requests.get()`, or heavy CPU computation inside a coroutine freezes *every* request in the process for its duration.

FastAPI's routing rule (Chapter 1, now in full): `async def` handlers run on the loop — they must `await` only non-blocking calls; plain `def` handlers are dispatched to a threadpool (~40 threads by default) — blocking code is safe there, at thread cost. The decision table:

Workload	Correct form
Async-capable I/O (<code>httpx</code> , <code>asyncpg</code> , <code>redis-py</code> <code>async</code>)	<code>async def</code> + <code>await</code>
Blocking-only library (<code>requests</code> , classic ORM, <code>boto3</code>)	plain <code>def</code> (threadpool)
Light CPU (< a few ms)	either
Heavy CPU (inference, parsing, crypto)	plain <code>def</code> , or offload: <code>run_in_executor</code> / process pool / worker queue

Table 8.1: Choosing `def` vs `async def` — the most important table in this book.

The deadliest production bug in FastAPI: `async def` handler → blocking call inside. It passes tests (single requests look fine), then under load p99 explodes for *all* endpoints at once. Detect it with a loop-lag monitor: a watchdog coroutine that sleeps 100 ms and alerts when actual wake-up time drifts.

8.2 Calling External APIs with httpx

External calls dominate modern API latency. Three non-negotiables: a *shared* client (connection pooling), *explicit timeouts* (the default “wait forever” is an outage), and *bounded retries with jitter* (only on idempotent operations):

```

1  import asyncio, random
2  import httpx
3  from contextlib import asynccontextmanager
4  from fastapi import FastAPI, HTTPException
5
6  @asynccontextmanager
7  async def lifespan(app: FastAPI):
8      app.state.http = httpx.AsyncClient(
9          base_url="https://api.partner.com",
10         timeout=httpx.Timeout(connect=2.0, read=5.0,
11                                write=5.0, pool=2.0),
12         limits=httpx.Limits(max_connections=100,
13                             max_keepalive_connections=20),
14     )
15     yield
16     await app.state.http.aclose()
17
18 app = FastAPI(lifespan=lifespan)
19
20 async def call_partner(client: httpx.AsyncClient, payload: dict,
21                       retries: int = 3) -> dict:
22     for attempt in range(retries):
23         try:
24             r = await client.post("/v1/score", json=payload)
25             if r.status_code in (502, 503, 504):
26                 raise httpx.HTTPStatusError("upstream", request=r.request,
27                                             response=r)
28             r.raise_for_status()
29             return r.json()
30         except (httpx.TimeoutException, httpx.HTTPStatusError):
31             if attempt == retries - 1:
32                 raise HTTPException(502, "Partner unavailable")
33             # exponential backoff with full jitter
34             await asyncio.sleep(random.uniform(0, 0.2 * 2 ** attempt))

```

Listing 8.1: A production httpx setup: lifespan-managed client, timeouts, retries.

Retry discipline: retry only transient classes (timeouts, 5xx, connection errors), never 4xx; cap total attempt time below *your* caller’s timeout; and add jitter — synchronized retries from many workers are a self-inflicted DDoS (the “thundering herd”). For repeated upstream failure, a circuit breaker stops calling a dead dependency entirely, failing fast until probes succeed.

8.3 BackgroundTasks vs Celery: Two Tiers of Deferred Work

BackgroundTasks runs callables *after the response is sent*, in the same process: perfect for fire-and-forget side effects (emails, audit logs, cache warming) where loss on process crash is acceptable.

```

1 from fastapi import BackgroundTasks
2
3 @app.post("/signup", status_code=201)
4 async def signup(payload: SignupIn, bg: BackgroundTasks, db: DB):
5     user = await user_service.create(db, payload)
6     bg.add_task(send_welcome_email, user.email) # after response
7     return {"id": user.id}

```

Listing 8.2: In-process background work — and its honest limits.

For anything that must *survive* — long-running jobs, retries, scheduled work, work that outlives deploys — you need a real queue: **Celery** (or RQ/arq/Dramatiq) with Redis or RabbitMQ as broker. The API enqueues a job, returns 202 Accepted with a job ID, and workers process independently; clients poll GET /jobs/{id} or receive a webhook on completion. The rule of thumb: *if losing the task on a pod restart is a bug, it does not belong in BackgroundTasks*.




Figure 8.1: The async-job pattern: accept fast, process reliably, report status — the backbone of every long-running-task API (and Chapter 9's batch inference).

8.4 Webhooks, Uploads, Streaming, and WebSockets

Webhooks (receiving): verify the HMAC signature *before* parsing (compute over the raw body with `hmac.compare_digest`), return 2xx immediately, and process via queue — senders time out fast and retry, so dedupe by event ID (idempotency again). **Webhooks (sending):** sign payloads, retry with backoff, and provide a redelivery dashboard — consumers *will* be down.

File uploads: `UploadFile` streams to a spooled temp file rather than loading into memory; copy in chunks to object storage (S3/GCS), never to local pod disk. Cap sizes at the proxy *and* in the handler; validate content type by sniffing magic bytes, not by trusting the client's header.

Streaming responses: `StreamingResponse` wraps an async generator; bytes flow as produced — the foundation for CSV exports, server-sent events, and LLM token streams (Chapter 9 builds on this):

```

1 from fastapi.responses import StreamingResponse
2
3 @app.get("/export/tasks.csv")
4 async def export_tasks(db: ADB):
5     async def rows():
6         yield "id,title,status\n"
7         async for t in repo.stream_all(db): # server-side cursor
8             yield f"{t.id},{csv_escape(t.title)},{t.status}\n"
9     return StreamingResponse(rows(), media_type="text/csv")

```

Listing 8.3: Streaming a large export without buffering it in memory.

WebSockets: `@app.websocket("/ws")` endpoints `accept()`, then loop on `receive_*/send_*` — a long-lived bidirectional channel for chat UIs and live dashboards. Authenticate during the handshake (token in query/header — cookies for browsers), expect disconnects (`WebSocketDisconnect`), and remember the operational shift: each connection is held state, so load balancers need long idle timeouts, deploys must drain connections, and horizontal scale-out needs a pub/sub backplane (Redis) to route messages across pods.

PRODUCTION CASE STUDY — The blocked event loop at a document-processing startup

A real and endlessly repeated incident shape: a team shipped an `async def /documents/analyze` endpoint that called a PDF-parsing library — synchronous, CPU-bound, 800 ms per document. Demos were flawless (one request at a time). At launch, 30 concurrent uploads made *every* endpoint — including `/health` — time out: the parser blocked the only event loop, so health checks queued behind PDF parsing; Kubernetes, seeing failed probes, restarted pods mid-request, multiplying retries into a feedback loop and a full outage. The three-line diagnosis came from a loop-lag metric added *after* the incident. Fixes applied, in order: change the handler to plain `def` (threadpool — immediate relief at thread-pool scale); then move parsing to Celery workers with a 202 + job-status API (correct architecture); keep `/health` trivially cheap and a dedicated loop-lag alert. Lessons: `async` is a contract, not a speed-up; health checks must not share fate with heavy work; and the job pattern (Figure 8.1) is the durable answer for >100 ms CPU work.

COST MODEL & METRICS — Concurrency capacity and the offload threshold

For an endpoint with CPU time T_{cpu} and I/O wait T_{io} per request, a single async worker saturates at

$$R_{max} \approx \frac{1}{T_{cpu}} \quad (\text{loop is busy only during CPU}), \quad \text{concurrency} \approx R_{max} \cdot (T_{cpu} + T_{io}).$$

With $T_{cpu} = 5$ ms, $T_{io} = 500$ ms: $R_{max} = 200$ req/s and ~ 101 requests in flight per worker. Threadpool capacity for blocking handlers: $R_{tp} = N_{threads} / (T_{cpu} + T_{io})$ — 40 threads at 505 ms ≈ 79 req/s, which is why blocking-only libraries cap throughput long before CPU does. Offload rule: move work to a queue when $T_{cpu} > 100$ ms or total time exceeds your gateway timeout. Queue sizing by Little's Law: backlog $L = \lambda \cdot W$; to keep wait W under 60 s at $\lambda = 50$ jobs/s with 2 s jobs, you need $\geq \lambda \cdot T_{job} = 100$ workers.

SYSTEM DESIGN SCENARIO — Design a video-thumbnail service: upload, process, notify

Requirements: clients upload videos up to 2 GB; thumbnails generated at multiple sizes (CPU/GPU, 10–60 s); client gets notified; 20K uploads/day with bursts; uploads must not tie up API capacity.

Reference approach: (1) Don't proxy bytes: API issues S3 presigned upload URLs (POST /videos → URL + video ID); client uploads directly to object storage — the API never touches 2 GB streams. (2) S3 event (or client confirm call) enqueues a processing job in Celery/SQS keyed by video ID; API returns 202 + /videos/{id}/status. (3) GPU worker pool sized by Little's Law against the 95th-percentile burst arrival rate; autoscale on queue depth, not CPU. (4) Workers write thumbnails to S3, update job state in Redis/Postgres, then notify: webhook to the customer (signed, retried, deduped by event ID) plus optional WebSocket push for browser clients. (5) Failure paths: per-job retry budget with dead-letter queue; partial results recorded; idempotent processing (same video ID reprocessed safely). (6) Observability: queue depth, job age p95, per-stage timing, DLQ rate — the four metrics that predict user-visible lag.

8.5 Interview Questions

Q1. Explain precisely what happens when an `async def` FastAPI handler makes a blocking call. Why does it degrade endpoints that share nothing with it?

Answer. Async handlers are coroutines scheduled on the process's single event loop. A blocking call (`requests.get`, `time.sleep`, heavy CPU) never yields control back to the loop, so for its duration *no other coroutine runs*: not other handlers, not keepalive handling, not health checks — they share the loop, which is the resource being monopolized. Under concurrency the effect compounds: k concurrent blocked requests serialize, queueing everything behind them; latency becomes $k \times T_{block}$ for innocent endpoints. Diagnosis: loop-lag monitoring (scheduled-vs-actual wake-up drift), py-spy dumps showing the loop thread in a sync call. Fixes: make the call truly async, switch the handler to `def` (threadpool), or offload via `run_in_executor/queue`. The senior framing: async is *cooperative* scheduling — one non-cooperator starves everyone.

Q2. Design a retry policy for calls to a flaky upstream. What can naive retries do to a degraded service, and what are the safeguards?

Answer. Naive retries amplify load exactly when the upstream is least able to absorb it: 3 retries turn 100% load into up to 400% during a brownout — often converting a slow service into a dead one (retry storm), and synchronized backoff across workers adds thundering-herd spikes. Policy: retry only transient errors (connect/read timeouts, 502/ 503/504, connection reset), never 4xx; only idempotent operations (or idempotency-keyed POSTs); exponential backoff *with full jitter* (`sleep(uniform(0, base·2k))`); a retry budget (e.g. retries ≤ 10% of requests — circuit-break beyond); total deadline below your own caller's timeout so retries don't outlive the request; and a circuit breaker that fails fast after threshold failures, probing periodically. Senior additions: honor `Retry-After`; propagate deadlines downstream (`X-Request-Deadline`) so the whole call tree respects one budget.

Q3. BackgroundTasks vs Celery: give the decision criteria and the failure modes of choosing wrong.

Answer. BackgroundTasks executes in-process after the response: zero infrastructure, but no persistence (pod crash/deploy loses it), no retries, no scheduling, no cross-process distribution, and it consumes the API process's own threadpool/loop — heavy background work degrades foreground latency. Celery (broker-backed) gives durability, acknowledged delivery, retries with backoff, scheduled/periodic jobs, independent worker scaling, and isolation. Criteria: tolerable-to-lose + short + light $\Rightarrow$ BackgroundTasks (audit log, cache warm, non-critical email); must-complete or long or heavy or bursty $\Rightarrow$ queue. Choosing wrong, direction 1: order-confirmation emails in BackgroundTasks silently vanish on every deploy — a trust-destroying, hard-to-detect bug. Direction 2: Celery for a 5 ms log write — operational overhead (broker, workers, monitoring) for nothing. Senior signal: name the middle option (arq/RQ for lighter-weight durable queues) and the transactional-outbox pattern when enqueue must be atomic with a DB commit.

Q4. How do you receive webhooks safely? Cover authenticity, replay, ordering, and the timeout contract.

Answer. Authenticity: verify an HMAC signature (e.g. X-Signature-256) computed over the raw request body with a shared secret, using `hmac.compare_digest` (constant-time); reject before any parsing. Replay: signatures should cover a timestamp header; reject stale ones (e.g. >5 min) and dedupe by event ID in Redis — senders deliver at-least-once, so duplicates are guaranteed, making handler idempotency mandatory. Ordering: never assume it — events arrive out of order under retries; design handlers around state-versioned upserts or sequence numbers, not event order. Timeout contract: senders time out in seconds and count slow responses as failures — so validate, persist/enqueue, return 200 immediately, and process asynchronously (the queue pattern). Senior additions: per-source rate limiting, a dead-letter store for unprocessable events with operator tooling, and monitoring delivery lag from the timestamp header.

Q5. When does WebSocket beat SSE/polling, and what operational complexity does it introduce?

Answer. WebSockets win when communication is genuinely bidirectional and low-latency: chat with typing indicators, collaborative editing, client-initiated subscriptions changing mid-connection. SSE (StreamingResponse of text/event-stream) wins for one-way server-push — notifications, progress, LLM token streams — being plain HTTP: proxies, auth, and reconnection (built-in Last-Event-ID) all just work. Polling remains right for low-frequency state checks. WebSocket operational cost: per-connection server state (memory, fd limits), LB configuration for long-lived upgraded connections, deploy draining (connections die on restart — clients need reconnect+resume protocols), horizontal scaling requires a pub/sub backplane (Redis) since connected clients are pinned to pods, and distinct auth (handshake-time, plus re-auth for long sessions). Senior rule: choose the weakest primitive that meets the requirement — most

“we need WebSockets” requirements are SSE-shaped.

Q6. Why must the shared `httpx.AsyncClient` live in the lifespan handler rather than being created per request?

Answer. Client construction is expensive: TLS context setup, and — the real cost — the connection pool starts empty, so per-request clients pay TCP+TLS handshakes (1–3 RTTs) on *every* upstream call, and leak connections if not closed. A lifespan-scoped client reuses keepalive connections (handshake amortized to near-zero), enforces global concurrency limits (`httpx.Limits`), and closes cleanly on shutdown. The same per-process-singleton logic applies to DB engines, model handles, and Redis clients — construct once in lifespan, inject via dependency.

Q7. What is the difference between `StreamingResponse` and building the full response in memory, and when is streaming mandatory?

Answer. A buffered response materializes the entire payload (memory $\propto$ response size; time-to-first-byte = full generation time), then sends. `StreamingResponse` pulls from a (async) generator and writes chunks as produced: constant memory, immediate first byte, and the connection drives backpressure. Mandatory when: payloads exceed sane memory (large exports), TTFB matters (LLM token streams — perceived latency is first-token latency), or the source is itself a stream (server-side DB cursor, upstream streaming API). Caveats: status code and headers are committed before generation — mid-stream errors can’t become a clean 500, so design in-band error signaling; and `Content-Length` is absent (chunked transfer).

Q8. How do you enforce an upload size limit robustly, and why is checking `Content-Length` insufficient?

Answer. `Content-Length` is a client claim — absent under chunked encoding and trivially falsifiable. Layered enforcement: (1) reverse proxy cap (`nginx client_max_body_size`) rejects the bulk cheaply before Python sees it; (2) in-handler enforcement that counts bytes while streaming (`await file.read(chunk)` in a loop, abort past the cap) so a lying client is cut off mid-stream rather than buffered; (3) for direct-to-S3 designs, presigned-POST policies carry `content-length-range` so storage enforces it. Also validate file *type* by magic bytes, not extension or claimed MIME, and scan before serving anything back to users.

Q9. Your Celery queue depth is growing without bound. Walk through the diagnosis and the levers.

Answer. Queue growth means arrival rate $\lambda >$ service rate μ . Diagnose which side moved: traffic spike ($\lambda \uparrow$ — check enqueue rate), slow tasks ($\mu \downarrow$ — per-task duration percentiles; a slow downstream dependency is the usual culprit), worker loss (crashes, OOM kills — worker count metrics), or poison messages (tasks failing + requeueing forever — retry counts, DLQ). Levers, by horizon: autoscale workers on queue depth (minutes); cap retries and dead-letter poison messages (immediately); shed load — reject or degrade low-priority enqueues at the API (429); split queues by priority so cheap critical tasks aren't stuck behind heavy ones; fix the slow dependency or add per-task timeouts. Prevent recurrence with an SLO on job age p95 and alerting on λ/μ ratio, not just absolute depth.

Q10. Explain backpressure in the context of an async API. Where does it come from and why is unbounded concurrency dangerous?

Answer. Backpressure is the mechanism by which a slower downstream limits a faster upstream's production rate. Async makes it easy to *accept* unbounded work — every request spawns a cheap coroutine — while downstreams (DB pool of 10, upstream API limits, worker memory) are finite: without limits, bursts become timeouts, pool exhaustion, and OOM kills rather than graceful queuing. Mechanisms: bounded pools (DB/httpx limits) naturally queue waiters; `asyncio.Semaphore` caps in-flight fan-out (`async with sem:`); queue maxsizes make producers wait; server-level concurrency limits (`uvicorn -limit-concurrency`) and gateway rate limits cap intake; and load shedding (fast 429/503 + `Retry-After` when saturated) protects latency for admitted requests. The design principle: every unbounded buffer in a system is a deferred outage — bound it and decide explicitly what happens at the bound.

FastAPI for AI, ML, RAG, and LLM Applications

This is the chapter the rest of the book has been building toward: serving ML models and LLM applications in production. We cover model loading and inference endpoint design, batching, LLM wrappers with token streaming, embedding APIs, vector databases and the full RAG architecture, guardrails and validation for AI inputs and outputs, model versioning, and the monitoring that AI systems — unlike ordinary APIs — cannot live without.

9.1 Serving ML Models: The Shape of the Problem

An ML inference service is an ordinary FastAPI app with three unusual properties: a **heavy startup artifact** (model weights take seconds to minutes to load — never per request), **CPU/GPU-bound hot paths** (everything Chapter 8 said about not blocking the loop applies doubly), and **probabilistic outputs** (correctness is a distribution, so monitoring must track *quality*, not just errors).

Load models once, in lifespan, and serve from memory:

```

1 from contextlib import asynccontextmanager
2 from anyio import to_thread
3 from fastapi import FastAPI
4 from pydantic import BaseModel, Field
5
6 @asynccontextmanager
7 async def lifespan(app: FastAPI):
8     app.state.model = load_model("models/sentiment-v3.onnx") # once
9     app.state.model_version = "sentiment-v3.2.1"
10    yield
11    app.state.model = None # release for clean shutdown
12
13 app = FastAPI(lifespan=lifespan, title="Inference API")
14
15 class PredictIn(BaseModel):
16     text: str = Field(min_length=1, max_length=4_000) # cost cap
17
18 class PredictOut(BaseModel):
19     label: str
20     confidence: float = Field(ge=0, le=1)
21     model_version: str # always identify the model

```

```

22
23 @app.post("/v1/predict", response_model=PredictOut)
24 async def predict(req: PredictIn):
25     # CPU-bound: run in threadpool, never on the event loop
26     label, conf = await to_thread.run_sync(
27         app.state.model.predict, req.text)
28     return PredictOut(label=label, confidence=conf,
29                       model_version=app.state.model_version)
30
31 @app.get("/v1/models") # version discovery endpoint
32 def model_info():
33     return {"version": app.state.model_version, "status": "ready"}

```

Listing 9.1: Model loading at startup; inference offloaded from the event loop.

Echoing `model_version` in every response is non-negotiable: when quality regresses, the first question is always “which model produced this?” — and during canary rollouts, multiple versions serve simultaneously.

9.1.1 Batch Inference

GPUs are throughput machines: predicting 64 inputs costs little more than one. Two batching modes: **offline** (client sends an array — cap its length; or a job-queue pipeline per Chapter 8 for large corpora) and **dynamic micro-batching** (the server holds requests for a few milliseconds, batches them through the GPU, and fans results back — the core trick of dedicated servers like Triton and vLLM). The trade is explicit: a ~5–10 ms added wait buys multiples of throughput.

9.2 Wrapping LLMs: Gateway Patterns and Streaming

Most LLM applications call hosted APIs (OpenAI, Anthropic) or a self-hosted engine (vLLM). Wrapping them in your own FastAPI service — an *LLM gateway* — is the pattern that pays for itself: one place for authentication, prompt templates, guardrails, caching, cost metering, provider failover, and model swaps without touching clients.

Streaming is what makes LLM UX tolerable: first tokens arrive in hundreds of milliseconds instead of the full completion’s tens of seconds. Server-sent events over `StreamingResponse`:

```

1 from fastapi.responses import StreamingResponse
2
3 class ChatIn(BaseModel):
4     message: str = Field(min_length=1, max_length=8_000)
5     max_tokens: int = Field(default=512, le=2_048) # cost ceiling
6
7 @app.post("/v1/chat/stream")
8 async def chat_stream(req: ChatIn, user: CurrentUser):
9     await guardrails.check_input(req.message) # pre-flight
10    async def tokens():
11        usage = {"in": 0, "out": 0}
12        try:
13            async for chunk in llm.stream(
14                system=PROMPT_TEMPLATE, user=req.message,
15                max_tokens=req.max_tokens,

```

```

16         ):
17             usage["out"] += 1
18             yield f"data: {chunk.json()}\n\n"          # SSE frame
19             yield "data: [DONE]\n\n"
20         finally:                                       # even on disconnect
21             await metering.record(user.id, usage)
22     return StreamingResponse(tokens(),
23                               media_type="text/event-stream")

```

Listing 9.2: An LLM gateway endpoint with SSE token streaming and guardrails.

Note the `finally:` clients disconnect mid-stream constantly, and cost metering must survive that. Note also `max_tokens` validated with a ceiling — a Pydantic constraint is literally a spending limit here.

9.3 Embeddings, Vector Databases, and RAG

Embedding APIs map text to vectors; design them batch-first (`texts: list[str]` with length caps) since embedding models are throughput-friendly, and version the model in responses — vectors from different models are *incomparable*, so a model upgrade means re-embedding the corpus.

RAG (Retrieval-Augmented Generation) grounds LLM answers in your documents: ingest (`chunk` → `embed` → store in a vector DB), then per query: `embed` the question, retrieve top-*k* similar chunks, optionally rerank, assemble a context-stuffed prompt, generate, and return the answer *with citations*.




Figure 9.1: RAG architecture: an ingestion pipeline (dashed) and a query pipeline (solid), with guardrails at both ends of the LLM.

Vector store choice: **pgvector** keeps vectors in Postgres — transactional consistency with your metadata, one database to operate, ideal up to millions of vectors; dedicated engines (Qdrant, Weaviate, Pinecone, Milvus) add scale and faster ANN at the cost of a second system. Start with pgvector; migrate behind the repository interface when scale demands.

9.4 Guardrails: Validating Inputs and Outputs of a Stochastic System

Input side: length and token caps (cost), prompt-injection screening (patterns + a small classifier on high-stakes routes), PII detection/redaction before text reaches third-party APIs, and topic policy checks. **Output side:** schema enforcement for structured outputs — parse the LLM’s JSON into a Pydantic model, and on `ValidationError` retry with the error message appended (one retry fixes most cases); citation verification for RAG (every cited chunk ID must exist in the retrieved set); content policy screening; and PII/secret scanning before anything reaches the client.

```

1 class ExtractedInvoice(BaseModel):
2     vendor: str
3     total: Decimal = Field(ge=0)
4     currency: Literal["USD", "EUR", "GBP"]
5     line_items: list[LineItem] = Field(max_length=200)
6
7 async def extract_invoice(text: str) -> ExtractedInvoice:
8     prompt = EXTRACT_PROMPT.format(text=text)
9     for attempt in range(2):
10         raw = await llm.complete(prompt, json_mode=True)
11         try:
12             return ExtractedInvoice.model_validate_json(raw)
13         except ValidationError as e:
14             prompt = f"{prompt}\n\nYour previous output failed "\
15                     f"validation:\n{e}\nReturn corrected JSON only."
16     raise HTTPException(502, "Model output failed validation")

```

Listing 9.3: Output validation: forcing a stochastic system through a typed contract.

This loop — generate, validate with Pydantic, feed errors back, retry — is the single most useful LLM-engineering pattern in this book: it converts a stochastic text generator into a typed API component with an explicit failure mode.

9.5 Versioning, Monitoring, and Humans in the Loop

Model versioning: model artifacts are immutable, semantically versioned deployables (registry: MLflow/W&B or object storage + manifest); APIs expose stable routes with the version in the *response*, canary new models on a traffic slice, and keep instant rollback (previous artifact still warm). Prompts deserve the same rigor: versioned templates in git, A/B-tested like models.

Monitoring adds two layers beyond Chapter 10’s standard stack: *AI performance* — time-to-first-token (TTFT), tokens/sec, token usage and cost per request/user/feature, cache hit rates — and *quality* — automated evals on sampled traffic (LLM-as-judge with calibration), retrieval relevance for RAG, output-validation failure rate, guardrail trigger rate, user feedback (thumbs up/down), and input distribution drift against the training/eval distribution.

Human-in-the-loop: below a confidence threshold (or on guardrail triggers), route outputs to a review API — a queue of pending items (GET /reviews/pending, POST /reviews/{id}/decision) feeding corrections back as training/eval data. Design the API so the human is a *component with an SLA*, not an afterthought: track reviewer latency and inter-reviewer agreement.

PRODUCTION CASE STUDY — Air Canada’s chatbot — the cost of unvalidated outputs

In 2024, a tribunal ordered Air Canada to honor a bereavement-fare policy its website chatbot had *invented*: the bot confidently described a refund policy that did not exist, a customer relied on it, and the court rejected the airline’s argument that the chatbot was “a separate legal entity responsible for its own actions.” The direct cost was small (\$812); the precedent and reputational cost were not: your model’s words are your company’s words. Engineering translation: the failure was architectural, not model-quality — there was no RAG grounding

in actual policy documents, no output validation requiring citations to canonical sources, no confidence-based routing to a human agent, and no fallback “I can’t answer that — here’s the policy page.” Every element was a known pattern (this chapter’s Figure 9.1 plus guardrails); the case is now the standard citation for why customer-facing LLM endpoints need them.

COST MODEL & METRICS — LLM serving economics

Per-request cost for a hosted LLM:

$$C_{req} = \frac{T_{in} \cdot p_{in} + T_{out} \cdot p_{out}}{10^6}$$

with T = tokens and p = price per million. At $p_{in} = \$3$, $p_{out} = \$15$, $T_{in} = 2,000$ (RAG context!), $T_{out} = 500$: $C_{req} = \$0.0135$ — so 1M requests/month = \$13,500, and *context size dominates*: RAG that stuffs 8K tokens of chunks quadruples input cost; retrieval precision is a cost lever, not just a quality lever. Levers, by typical ROI: response caching for repeated queries (hit rate h scales cost by $1 - h$), prompt/context compression, output caps (`max_tokens`), model routing (cheap model for easy queries, escalate hard ones), and batch APIs (often 50% discounts) for offline work. Self-hosting crossover: a \$2/hr GPU serving R req/s beats the API when $C_{req} \cdot R \cdot 3600 > 2$, i.e. $R \gtrsim 0.04$ req/s sustained at the prices above — but add the engineering-time cost honestly. Always meter per user/feature: one runaway agent loop can spend a month’s budget in a day.

SYSTEM DESIGN SCENARIO — Design a document-Q&A (RAG) API for an enterprise knowledge base

Requirements: 200K internal documents with per-team access control, answers must cite sources, p95 first-token < 1.5 s, 50 concurrent users, hallucination is a firing-offense word in the RFP.

Reference approach: (1) Ingestion pipeline (queue-based, Chapter 8): parse → chunk (~500 tokens, overlap) → embed (batched) → pgvector, with document ACL tags stored on every chunk. (2) Query path: validate + guardrail input → embed query → ANN top-20 *filtered by the caller’s ACL in the SQL query* (security boundary in the database, not the prompt — never let the LLM see chunks the user can’t) → rerank to top-5 → prompt with instruction to answer *only* from context and cite chunk IDs → stream via SSE. (3) Output validation: every cited ID must be in the retrieved set; answers without citations are replaced by an honest “not found in the knowledge base” — refusal is a feature. (4) Quality loop: log query/chunks/answer (with retention policy); nightly eval set scored by LLM-judge; retrieval-relevance and citation-failure dashboards; thumbs-down routes to human review. (5) Capacity: embedding and LLM calls are I/O — async handlers, shared clients, semaphore-bounded fan-out; pgvector HNSW index sized for $200K \times$ chunks-per-doc vectors easily on one Postgres.

9.6 Interview Questions

Q1. Why must ML models load in the lifespan handler, and what goes wrong with lazy per-request or import-time loading?

Answer. Per-request loading is catastrophic: seconds of disk/GPU transfer per call, memory churn, and concurrent requests each loading copies until OOM. Import-time loading works but couples loading to module import — tests, linters, and CLI tools importing the app pay the cost (or crash without GPU), and there is no clean shutdown hook. Lifespan loading is correct: it runs once per worker before traffic, plays well with health/readiness probes (pod reports ready only after the model is warm — critical for rolling deploys), supports clean release on shutdown, and TestClient's context manager exercises it. Senior additions: with N workers you hold N model copies — size worker count by memory, or use a dedicated inference server (Triton, vLLM) behind FastAPI when models are large; and lazy-load *rarely used* model variants behind a readiness flag rather than blocking startup for all of them.

Q2. Design the streaming architecture for an LLM chat endpoint end to end — protocol choice, disconnects, metering, and what changes at the infrastructure layer.

Answer. Protocol: SSE over StreamingResponse for one-way token flow (WebSocket only if mid-generation client→server interaction is needed); frames are data: JSON chunks with a terminal [DONE]. Handler: async generator awaiting the provider's stream; guardrails run pre-flight on input; output moderation runs incrementally (buffer sentence-windows) or post-hoc since you cannot unsay streamed tokens — for strict-policy products, stream from a buffered-and-checked window with a small delay. Disconnects: ubiquitous — wrap the generator in try/finally so token metering, logging, and provider stream cancellation always run; propagate cancellation upstream to stop paying for abandoned generations. Infrastructure: disable proxy buffering (nginx X-Accel-Buffering: no), raise idle timeouts on LB/gateway, and monitor TTFT separately from total latency — TTFT is the UX metric. Capacity: each active stream holds a slot for tens of seconds, so concurrency limits and per-user stream caps prevent a few chatty clients from starving the service.

Q3. A RAG system gives confidently wrong answers. Walk through your diagnosis across the pipeline.

Answer. Decompose — RAG failures are almost always retrieval failures wearing a generation costume. (1) *Retrieval*: for failing queries, inspect the retrieved chunks — were the right documents present? If not: chunking too coarse/fine, embedding model mismatch with domain vocabulary, missing metadata filters, or the answer simply isn't in the corpus (then the correct output is refusal). Measure recall@k on a labeled eval set. (2) *Ranking*: right doc retrieved but ranked below the context cutoff ⇒ add a reranker. (3) *Prompt assembly*: context truncated, poorly ordered (relevant chunk buried middle — position effects are real), or instructions

weak (“answer only from the context, cite sources, else say you don’t know”). (4) *Generation*: model ignoring context $\Rightarrow$ stronger grounding instructions, citation-required validation, smaller context with higher precision. (5) *Systemic*: no citation verification, no refusal path, no eval harness — the meta-fix is a nightly evaluation suite so regressions surface as metrics, not incidents. Senior signal: you measure each stage (recall@k, rerank lift, citation-failure rate) instead of prompt-tweaking blindly.

Q4. How do you make a stochastic LLM safe to use as a typed component in a larger system?

Answer. Contract it: (1) request structured output (JSON mode / function-calling APIs); (2) parse into a Pydantic model — the same validation boundary as any untrusted input, with field constraints encoding business rules (`total ≥ 0`, enum currencies); (3) on `ValidationError`, retry once with the error text appended — empirically fixes the majority; (4) on second failure, raise a typed error to a defined fallback: human review queue, cheaper deterministic extractor, or explicit 502 — never silently pass malformed data downstream; (5) log validation-failure rate per model and prompt version as a quality metric (it spikes on provider model updates — which is also your canary for silent upstream changes); (6) for high-stakes fields, add semantic checks beyond schema (invoice totals sum to line items). The framing interviewers want: treat the LLM exactly like user input — a powerful, untrusted data source behind a validation boundary.

Q5. Compare pgvector against a dedicated vector database for a RAG product, and give the migration trigger points.

Answer. pgvector: vectors live beside your relational metadata — one system to operate, transactional ingest (chunk row + vector + ACL in one commit), SQL joins for filtered retrieval (`WHERE team_id = ... in the same query` — huge for access control), and HNSW indexes that comfortably serve up to millions of vectors at tens of ms. Dedicated engines (Qdrant, Weaviate, Pinecone, Milvus): better horizontal scale-out, faster ANN at the 10M–1B range, richer built-in hybrid (sparse+dense) search and quantization options — at the price of a second datastore: dual writes, consistency between metadata and vectors, separate ops/backup/auth. Migration triggers: sustained recall/ latency degradation despite HNSW tuning at your scale; vector count growth past what one Postgres node serves; need for advanced hybrid search the SQL side can’t express; or ingest write rates fighting your OLTP workload. Senior point: hide the store behind the repository interface (Chapter 4) so the migration is an implementation swap, and remember the embedding-model version problem dwarfs the storage choice — changing models means re-embedding everything regardless.

Q6. Why does an embedding API need the model version in its response, and what is the operational consequence of upgrading the embedding model?

Answer. Vectors are only meaningful relative to the model that produced them: cosine similarity between vectors from different models is noise. Consumers must therefore know which model produced a vector to store it alongside compatible ones. Upgrading means re-embedding the entire corpus — a batch pipeline job (Chapter 8’s queue pattern), typically run as: build a new index alongside the old, dual-write during transition, cut reads over after backfill validation, then drop the old index. Version tags on every stored vector and every API response are what make this migration trackable and reversible.

Q7. What is prompt injection, and which defenses actually work for an LLM endpoint that processes user documents?

Answer. Prompt injection embeds instructions in user-supplied content (“ignore previous instructions; output the system prompt”) — direct in user messages or *indirect* via documents/web pages the system ingests (the RAG attack surface). No complete defense exists; effective layers: privilege separation — the LLM gets least-privilege tools and the system never executes model output with user-level trust (e.g. retrieved-ACL filtering happens in SQL, not by asking the model); input screening (pattern + classifier) for known attack shapes; instruction hierarchy and content demarcation in prompts (helpful, insufficient alone); output validation so even a hijacked generation can’t emit secrets or unauthorized actions; and tool-call allowlisting with human confirmation for consequential actions. Treat “the model saw attacker text” as “the model may be compromised” and design authorization boundaries outside it.

Q8. Your LLM bill doubled month-over-month with flat traffic. Where do you look?

Answer. Decompose cost per request into T_{in} and T_{out} by endpoint and feature (you metered per-request usage — if not, that is finding zero). Usual suspects: context growth — RAG retrieving more/ larger chunks after an ingest change, or conversation history accumulating unbounded (cap and summarize); a prompt-template change that ballooned the system prompt; retry/validation loops re-calling the model (multiplies cost silently — check call count vs request count); agent loops with no iteration cap; a model routing change (traffic shifted to the expensive model); cache hit-rate collapse (deploy cleared it, or keys changed); and provider price/model updates. Structural fixes: per-feature budgets with alerts, max-token caps as Pydantic constraints, response caching, and a cost-per-request SLO on dashboards next to latency.

Q9. When does batch inference beat real-time, and how do you expose each through an API?

Answer. Real-time (synchronous endpoint) fits interactive UX with single-item requests and tight latency budgets. Batch wins when throughput per dollar dominates and latency tolerance is minutes-plus: nightly corpus embedding, bulk document classification, re-scoring — GPUs amortize fixed per-call overhead across the batch, and hosted batch APIs are often half price. API shapes: real-time = `POST /v1/predict` with one item (optionally a small array, capped); batch = the job pattern — `POST /v1/batches` with a file/dataset reference returning 202 + job ID, status polling, results to object storage, webhook on completion. The middle ground is server-side dynamic micro-batching: clients see single-item real-time semantics while the server batches concurrent requests within a few-ms window — the best of both when concurrency is high.

Q10. Design the human-in-the-loop component for an AI document-extraction service. What makes it an engineering problem rather than a UI problem?

Answer. Routing: every extraction carries confidence (model scores + validation outcomes); below threshold — or on guardrail/validation failures — items enter a review queue instead of the response path, with the API returning “pending review” status (the job pattern, so clients handle async outcomes from day one). The review API: list pending items (with model output prefilled for correction — review is editing, not re-doing), claim/lock semantics so two reviewers don’t collide, decision endpoints recording corrected output + reason codes, full audit trail. Engineering substance: the queue has SLAs (reviewer latency is part of your product latency for routed items — monitor it), thresholds are tunable trade-offs (precision of automation vs review cost: $\text{route rate} \times \text{volume} \times \text{minutes-per-review} = \text{headcount}$), corrections feed evaluation and retraining datasets (closing the data flywheel), and inter-reviewer agreement is measured to detect ambiguous taxonomy. Done right, the threshold becomes a dial: quality, cost, and latency become an explicit, tunable triangle.

Production Deployment, Scaling, and Observability

Everything before this chapter made the service correct; this chapter makes it survive contact with production: servers and workers, Docker, configuration, cloud deployment and Kubernetes, CI/CD, health checks, logging, metrics, tracing, caching, performance tuning, load testing — and the production checklist that ties the whole book together.

10.1 Servers and Workers: Uvicorn and Gunicorn

Uvicorn is the ASGI server; one Uvicorn process is one event loop on one core. Production needs one worker per core plus supervision. Two standard arrangements: `uvicorn -workers N` (built-in multi-process, fine on Kubernetes where the platform supervises pods) or **Gunicorn with Uvicorn workers** (`gunicorn -k uvicorn.workers.UvicornWorker -w N`) — Gunicorn adds mature process management: worker restarts on crash, `max-requests` recycling against slow leaks, graceful timeouts. On Kubernetes, a common simplification is *one worker per pod* and horizontal scaling by pod count: the scheduler does the supervision, and resource accounting stays clean. Worker math: I/O-bound async apps, $N \approx \text{cores}$; CPU-heavy apps, $N = \text{cores}$ with the work offloaded (Chapter 8); memory bounds it all — N workers = N model copies (Chapter 9).

10.2 Dockerizing FastAPI

```

1 FROM python:3.12-slim AS builder
2 WORKDIR /app
3 COPY requirements.txt .
4 RUN pip install --no-cache-dir --prefix=/install -r requirements.txt
5
6 FROM python:3.12-slim
7 RUN useradd --create-home appuser
8 WORKDIR /app
9 COPY --from=builder /install /usr/local
10 COPY app/ app/
11 COPY alembic.ini .
12 COPY migrations/ migrations/
13 USER appuser
14 EXPOSE 8000

```

```

15 HEALTHCHECK --interval=30s --timeout=3s \
16   CMD python -c "import httpx;
    ↪ httpx.get('http://localhost:8000/health/live').raise_for_status()"
17 CMD ["uvicorn", "app.main:app", "--host", "0.0.0.0", "--port", "8000"]

```

Listing 10.1: A production Dockerfile: multi-stage, non-root, healthcheck.

The essentials: multi-stage build (no compilers in the runtime image), slim base, non-root user, pinned dependencies (lock file), and `.dockerignore` excluding `.env`, `.git`, tests. For local parity, Docker Compose runs the full topology — app, Postgres, Redis, worker — so “works on my machine” means the machine that matters:

```

1  services:
2    api:
3      build: .
4      ports: ["8000:8000"]
5      environment:
6        DATABASE_URL: postgresql+psycopg://app:secret@db:5432/app
7        REDIS_URL: redis://cache:6379/0
8      depends_on: [db, cache]
9    worker:
10     build: .
11     command: celery -A app.worker worker --concurrency=4
12     environment: {REDIS_URL: "redis://cache:6379/0"}
13     depends_on: [cache]
14    db:
15     image: postgres:16
16     environment: {POSTGRES_USER: app, POSTGRES_PASSWORD: secret,
17                  POSTGRES_DB: app}
18     volumes: [pgdata:/var/lib/postgresql/data]
19    cache:
20     image: redis:7
21  volumes:
22    pgdata: {}

```

Listing 10.2: docker-compose.yml: the production topology on a laptop.

10.3 Configuration, Cloud Targets, and Kubernetes

Environment-based configuration via `pydantic-settings`: a `Settings` class reading env vars (12-factor), validated at startup — a missing `DATABASE_URL` should crash the boot, not the first request. Secrets come from the platform’s secret store; the same image runs in every environment with different env.

Cloud targets, in increasing control/complexity: serverless containers (**Cloud Run**, App Runner, Container Apps — scale-to-zero, per-request billing; watch cold starts for model-heavy apps), managed platforms (ECS/Fargate), and **Kubernetes** when you need its ecosystem. The Kubernetes vocabulary that matters for an API: a *Deployment* (replica management, rolling updates), a *Service* + *Ingress* (routing), an *HPA* (autoscaling on CPU or custom metrics like queue depth), resource *requests/limits* (capacity math), and the two probes:

```

1 @app.get("/health/live")           # "is the process stuck?" -> restart me
2 def live():
3     return {"status": "ok"}        # no dependencies: process responds
4
5 @app.get("/health/ready")          # "can I serve?" -> route traffic to me
6 async def ready(db: ADB):
7     await db.execute(text("SELECT 1"))    # critical deps only
8     if not app.state.model:              # Chapter 9: model warm?
9         raise HTTPException(503, "model not loaded")
10    return {"status": "ready"}

```

Listing 10.3: Liveness vs readiness: two different questions.

Getting these backwards is a classic outage: a liveness probe that checks the database restarts every pod when the DB blips — turning a degradation into a full outage. Liveness = process health only; readiness = dependency health.

CI/CD pipeline: lint + types → tests (Chapter 7) → build image (tagged by git SHA) → scan → push → run Alembic migrations (expand/contract, Chapter 5) → progressive rollout (rolling or canary) → automated rollback on health/metric regression.

10.4 Observability: Logs, Metrics, Traces

The three pillars answer different questions. **Structured logs** (JSON, one event per line: timestamp, level, message, request ID, user ID, latency) answer “what happened in this specific case?” **Metrics** (Prometheus) answer “how is the system behaving in aggregate?” — the RED trio per endpoint: Rate, Errors, Duration histograms, plus the resource gauges this book has flagged (pool checkout wait, queue depth, loop lag, token costs). **Distributed traces** (OpenTelemetry) answer “where did this request spend its time across services?” — indispensable once a request touches gateway → API → DB → LLM provider. opentelemetry-instrument auto-instruments FastAPI, httpx, and SQLAlchemy; propagate the trace ID into every log line so the pillars cross-link. Grafana sits on top: one dashboard per service with RED panels, resource panels, and SLO burn rate.




Figure 10.1: The production topology: replicated API pods behind an ingress, shared state in Postgres/Redis, workers for deferred jobs, and an observability plane receiving logs, metrics, and traces from everything.

10.5 Caching, Performance Tuning, and Load Testing

Redis caching, in order of safety: idempotent expensive reads (cache-aside with TTL: check Redis, on miss compute and SETEX), session/rate-limit state (Chapter 6), and LLM response caching keyed on normalized prompt + model version (Chapter 9 — often the largest single cost saver). The two hard problems are honest TTL choice (staleness budget is a product decision) and invalidation: prefer key-versioning (`user:42:v{version}`) over scanning, and stampede protection (per-key locks or probabilistic early refresh) for hot keys.

Performance tuning, in the order that actually finds wins: measure first (profile under realistic load — py-spy, OTel spans); fix N+1s and missing indexes (Chapter 5 — the majority of real-world API latency); eliminate event-loop blocking (Chapter 8); cache what is expensive and repeatable; tune serialization only when measured (ORJSON response class for large payloads); and scale horizontally last, because scaling a service with a blocked loop or N+1 queries just multiplies the waste.

Load testing with Locust or k6: model realistic scenarios (mixed endpoints, think time, data variety — not one URL in a loop), ramp to find the knee of the latency curve, run soak tests (hours, for leaks and pool drift) and spike tests (cold-cache behavior). Test *through* the real ingress path, against production-shaped data volumes, with observability on — the load test is also a test of your dashboards.

COST MODEL & METRICS — Capacity, autoscaling, and the cost of latency

Capacity per pod from load testing: R_{pod} = req/s at the latency knee. Pods needed at peak with utilization target u (headroom for spikes and rolling deploys):

$$N_{pods} = \left\lceil \frac{R_{peak}}{u \cdot R_{pod}} \right\rceil, \quad u \approx 0.6\text{--}0.7.$$

Example: 2,000 req/s peak, R_{pod} = 180 req/s, u = 0.65 $\Rightarrow$ 18 pods; at 1 vCPU/2GB ($\sim$ \$35/month each) $\approx$ \$630/month — before the database, which usually costs more. Autoscaling on CPU suffices for I/O-bound APIs; queue-depth or RPS-based scaling reacts faster for bursty traffic. And the business side: industry studies (Amazon, Google) repeatedly find conversion drops of $\sim$ 1% per 100 ms of added latency — the p99, not the average, is where users live. Track cost per million requests as a first-class metric next to latency; optimization without that denominator is guesswork.

PRODUCTION CASE STUDY — The readiness-probe cascade — a composite of real K8s postmortems

A pattern documented in numerous public Kubernetes postmortems (and likely running somewhere right now): a team points both liveness and readiness probes at a health endpoint that checks the database. A 30-second database failover begins; every pod's probes fail simultaneously. Readiness failure correctly removes pods from the load balancer — but liveness failure *restarts* them all. Restarted pods come up cold (Chapter 9: model loading takes 90s), fail readiness during warm-up, and the HPA — seeing CPU near zero because no traffic flows — scales *down*. The DB recovers in 30 seconds; the service takes 11 minutes to recover, because the platform spent the incident killing healthy-but-waiting processes. Fixes: liveness checks

process health only (never dependencies); readiness checks critical dependencies with a failure threshold tolerant of blips; startup probes cover long warm-ups so liveness doesn't kill booting pods; and PodDisruptionBudgets + pre-stop drain hooks protect capacity during rollouts. The meta-lesson: self-healing automation amplifies misconfigured signals — the platform did exactly what the probes told it to.

SYSTEM DESIGN SCENARIO — Take the Chapter-9 RAG API from one VM to a real production platform

Requirements: the document-Q&A service now serves 3 products, 99.9% availability target, p95 first-token < 1.5s, deploys daily, 2-person platform budget.

Reference approach: (1) *Platform:* managed Kubernetes or Cloud Run — with a 2-person team, buy the control plane. Separate deployments for API pods (async, lightweight) and ingestion workers (queue consumers); embedding/LLM calls are external, so API pods scale on RPS. (2) *Data:* managed Postgres (with pgvector) + managed Redis; PgBouncer; automated backups with restore drills. (3) *Config/secrets:* pydantic-settings from env; secrets in the platform store; one image, three environments. (4) *Pipeline:* GitHub Actions — lint/type/test gates, image build by SHA, migration job, canary 10% for 15 minutes gated on error rate and TTFT, then full rollout; rollback is redeploying the previous SHA. (5) *Observability:* RED metrics + TTFT, token cost/request, retrieval recall sampling; traces across API → vector DB → LLM; SLO 99.9% with burn-rate alerts paging only on user-impacting symptoms. (6) *Resilience:* per-dependency timeouts shorter than the gateway timeout, circuit breaker on the LLM provider with a degraded “retrieval-only” fallback response, load-shedding 429s above measured capacity, PodDisruptionBudgets, and a quarterly game-day exercising the DB-failover and provider-outage runbooks.

10.6 The Production Checklist

Before the first real user: every endpoint authenticated and authorized (deny by default) with rate limits on auth and expensive routes; secrets in a manager, never in code or images; configuration validated at boot; migrations reviewed and reversible; pools sized by arithmetic (Chapter 5); timeouts on *every* external call, shorter than your own; structured logs with request IDs, RED metrics, traces, and alerts on symptoms (SLO burn) not causes; liveness/readiness/startup probes correctly separated; graceful shutdown draining in-flight requests; load test at the knee + soak test passed; backups restorable (tested, not assumed); dependency and image scanning in CI; runbooks for the top five failure modes; and a rollback path exercised at least once on purpose.

10.7 Interview Questions

Q1. Why do liveness and readiness probes need different implementations, and what failure mode does conflating them create?

Answer. They answer different questions with different remedies. Liveness: “is this process broken beyond self-recovery?” — remedy is restart; it must check only process-local health (event loop responds), because restarting a pod does not fix a down database. Readiness: “should this pod receive traffic right now?” — remedy is removal from the endpoint set; it should check critical dependencies and warm-up state (model loaded). Conflated probes that both check dependencies turn any shared-dependency blip into a mass-restart cascade: all pods restart simultaneously, lose warm caches/models, fail readiness during boot, and the platform amplifies a 30-second degradation into a multi-minute outage. Senior additions: startup probes for long boot sequences (so liveness doesn’t kill loading pods), failure thresholds tolerant of transient blips, and the principle that self-healing automation acts on your signals — wrong signals, wrong healing.

Q2. Design a zero-downtime deploy for a FastAPI service with database migrations and 30-second-lived in-flight requests.

Answer. Sequence: (1) migrations first, expand/contract style (Chapter 5) — new schema must serve *both* old and new code, applied as a pre-rollout job with a lock/lease so only one runner executes. (2) Rolling update with `maxUnavailable=0`, `maxSurge=1`-style settings: new pods join only after readiness passes (dependencies up, model warm). (3) Graceful termination: on SIGTERM the pod fails readiness (stops new traffic), a `preStop` sleep covers endpoint-propagation lag, Uvicorn finishes in-flight requests within `terminationGracePeriodSeconds > 30s + buffer`; WebSocket/SSE connections need client reconnect logic since they will be cut. (4) Gate progression on canary metrics (error rate, p95) with automated rollback to the previous image SHA — which is why images are SHA-tagged and the previous schema still works (contract step deferred until the new code is proven). Senior signals: connection draining at the LB layer too, `PodDisruptionBudgets` against simultaneous voluntary evictions, and rehearsed rollback.

Q3. Your p99 latency tripled but p50 is unchanged. Walk through your diagnosis using the three observability pillars.

Answer. p50-stable/p99-degraded means a *subset* of requests hits a slow path — the median request is fine. Metrics first: slice the duration histogram by endpoint, status, pod, and time — is the tail one endpoint (code path), one pod (bad node, GC, noisy neighbor), or global-periodic (cron, cache expiry stampede)? Then traces: pull exemplar traces from the slow bucket and read where the time went — DB span (pool checkout wait? lock contention? a

query whose plan flipped?), external call (retries? upstream tail?), or gap-between-spans (event-loop blocking — Chapter 8’s signature, confirmed by loop-lag metrics). Logs: correlate the slow request IDs — shared user? large payloads? specific tenant whose data shape triggers N+1? Common culprits, in base-rate order: pool exhaustion under a slow query, periodic cache misses/stampedes, retry amplification on a flaky upstream, GC/threadpool saturation, and one heavy tenant. The senior marker is the method: histogram → exemplar trace → correlated logs, not guessing.

Q4. Cache-aside with TTL: walk through the stampede problem and its solutions, and explain why invalidation is harder than caching.

Answer. Stampede: a hot key expires; 500 concurrent requests all miss, all recompute the expensive value (or hammer the DB), latency spikes, and under load the recompute storm can take the backend down — the cache was load-bearing. Solutions: per-key mutex/lease (one request recomputes, others serve stale or wait), probabilistic early refresh (each request refreshes early with probability rising near expiry — no coordination), background refresh for known-hot keys, and jittered TTLs so cohorts don’t expire together. Invalidation is harder because it requires knowing every write path that affects a cached value: TTL bounds staleness but doesn’t eliminate it; explicit invalidation must be wired into each mutation (and misses one eventually); key-versioning (bump `user:42:v7` on write) sidesteps deletion but moves the problem to version storage. Senior framing: declare a staleness budget per data class (product decision), choose the cheapest mechanism meeting it, and monitor hit rate + staleness, since a silent hit-rate collapse is a latency *and* cost incident (Chapter 9’s LLM cache doubly so).

Q5. How would you load test this book’s RAG service meaningfully? What do naive load tests get wrong?

Answer. Naive tests hammer one endpoint with identical payloads, no auth, against staging with toy data — measuring the cache and the happy path. Meaningful: (1) scenario mix from production traffic ratios (queries, ingestion, status polls) with think times; (2) realistic data variety — distinct queries (LLM/embedding caches behave honestly), production-scale vector corpus (ANN latency depends on index size), realistic ACL distribution; (3) through the full ingress path with real auth; (4) streaming-aware metrics: TTFT and inter-token latency, not just request duration — a tool must consume SSE properly; (5) ramp to find the knee, then soak (hours — pool drift, memory leaks) and spike (cold caches, autoscaler reaction time); (6) cost telemetry on: a load test against paid LLM APIs needs a fake or a budget. Output is not “it survived”: it is R_{pod} at the knee, the bottleneck’s identity (loop? pool? provider rate limit?), and autoscaler behavior — the inputs to the capacity formula.

Q6. Gunicorn with Uvicorn workers vs plain Uvicorn -workers vs one-process-per-pod on Kubernetes: how do you choose?

Answer. They differ in who supervises processes. Gunicorn+UvicornWorker: battle-tested supervision — crash restarts, `max_requests` recycling (leak mitigation), graceful worker time-outs; right for VMs/ bare metal where nothing else supervises. Plain `uvicorn -workers N`: lighter, fewer layers. One worker per pod on K8s: the kubelet and Deployment controller *are* the supervisor — restarts, scaling, and resource accounting per process; simplest mental model, cleanest metrics (one loop per pod), at the cost of more pods. Decision: K8s $\Rightarrow$ usually 1 worker/pod (or a few to amortize sidecar overhead); single VM $\Rightarrow$ Gunicorn; serverless containers $\Rightarrow$ the platform decides concurrency. Always: workers $\times$ pods sized against DB pool math (Chapter 5).

Q7. What belongs in structured logs for an API, and what must never appear?

Answer. Belongs: timestamp, level, event name, request ID (and trace ID), method/route (the *template*, not the raw URL), status, duration, user/tenant ID (as opaque IDs), and bounded business context (order ID, job ID). Never: passwords, tokens/API keys, full Authorization headers, raw request bodies containing PII, secrets in connection strings, and full prompt/response text in LLM systems without an explicit retention/consent policy (log a content hash + token counts instead). Enforcement is structural: `SecretStr`, serializers with redaction filters, and log-scanning canaries — because one leaked credential in logs has the blast radius of a breach (logs are widely readable and long-retained).

Q8. Explain SLO-based alerting and why paging on “CPU above 80%” is an anti-pattern.

Answer. An SLO sets a user-facing target (99.9% of requests succeed under 500 ms over 30 days), implying an error budget (0.1%). Alerts fire on *budget burn rate* — fast burn (e.g. consuming 2% of budget in an hour) pages, slow burn opens tickets. CPU at 80% is a *cause* signal: it may be perfectly healthy (good utilization!) or already-too-late, it pages humans for conditions users don’t feel (alert fatigue), and it misses user pain with idle CPU (blocked loop, upstream failures). Symptoms (error rate, latency SLI) page; causes (CPU, pool, queue depth) appear on the dashboard you open *after* the page. The discipline: every page should mean “a user is being hurt and a human must act now.”

Q9. Your service must survive its primary LLM provider's outage. What does graceful degradation look like?

Answer. Layers: (1) detection — circuit breaker on the provider client trips after threshold failures, failing fast instead of queueing 30s timeouts. (2) Fallback chain: secondary provider (kept warm with a trickle of traffic, abstraction layer from Chapter 9's gateway pattern), a smaller self-hosted model for degraded-but-present answers, or feature-specific fallbacks — RAG can return retrieval-only results ("here are the relevant documents") which preserves most user value at zero LLM cost. (3) Honest communication: degraded responses are labeled; status page reflects partial outage. (4) Protection: load-shed non-critical LLM features first (summaries off, core Q&A on) via feature flags. (5) Verification: provider-outage game days — an untested fallback is a hypothesis, not a capability. The principle: availability of *the product* is not binary in the availability of its dependencies.

Q10. A pod is OOM-killed every few hours under steady traffic. Walk through the investigation.

Answer. Confirm the kill (`kubectl describe pod`: OOMKilled, exit 137) and pull memory metrics: sawtooth-to-limit indicates a leak or accumulation; instant spikes indicate a single huge allocation (unbounded request body, giant query result, batch too large). Leak hunt: heap profiling (`tracemalloc`/`memray`) under replayed traffic; usual suspects — unbounded in-process caches (`@lru_cache` on unbounded keys, dicts keyed by user input), growing global lists/registries, response buffering instead of streaming (Chapter 8), sessions/clients created per request and never closed, and ML tokenizer/model copies per worker (Chapter 9). Mitigations while fixing: right-size limits (the limit may simply be below legitimate working set — check RSS at healthy steady state), Gunicorn `max_requests` recycling as a tourniquet, and bound every buffer (upload caps, query LIMITs, cache max sizes). Verify the fix with a soak test (Chapter-10 load testing) watching RSS flatline.